

A Comprehensive Survey on Graph Anomaly Detection With Deep Learning

Xiaoxiao Ma, Jia Wu^{ID}, Senior Member, IEEE, Shan Xue^{ID}, Jian Yang^{ID}, Chuan Zhou^{ID}, Quan Z. Sheng^{ID}, Hui Xiong^{ID}, Fellow, IEEE, and Leman Akoglu^{ID}

Abstract—Anomalies are rare observations (e.g., data records or events) that deviate significantly from the others in the sample. Over the past few decades, research on anomaly mining has received increasing interests due to the implications of these occurrences in a wide range of disciplines - for instance, security, finance, and medicine. For this reason, anomaly detection, which aims to identify these rare observations, has become one of the most vital tasks in the world and has shown its power in preventing detrimental events, such as financial fraud, network intrusions, and social spam. The detection task is typically solved by identifying outlying data points in the feature space, which, inherently, overlooks the relational information in real-world data. At the same time, graphs have been prevalently used to represent the structural/relational information, which raises the *graph anomaly detection problem* - identifying anomalous graph objects (i.e., nodes, edges and sub-graphs) in a single graph, or anomalous graphs in a set/database of graphs. Conventional anomaly detection techniques cannot tackle this problem well because of the complexity of graph data (e.g., irregular structures, relational dependencies, node/edge types/attributes/directions/multiplicities/weights, large scale, etc.). However, thanks to the advent of deep learning in breaking these limitations, graph anomaly detection with deep learning has received a growing attention recently. In this survey, we aim to provide a systematic and comprehensive review of the contemporary deep learning techniques for graph anomaly detection. Specifically, we provide a taxonomy that follows a task-driven strategy and categorizes existing work according to the anomalous graph objects that they can detect. We especially focus on the challenges in this research area and discuss the key intuitions, technical details as well as relative strengths and weaknesses of various techniques in each category. From the survey results, we highlight 12 future research directions spanning unsolved and emerging problems introduced by graph data, anomaly detection, deep learning and real-world applications. Additionally, to provide a wealth of useful resources for future studies, we have compiled a set of open-source implementations, public datasets, and commonly-used evaluation metrics. With this survey, our goal is to create a “one-stop-shop” that provides a unified understanding of the problem categories and existing approaches, publicly available hands-on resources, and high-impact open challenges for graph anomaly detection using deep learning.

Index Terms—Anomaly detection, outlier detection, fraud detection, rumor detection, fake news detection, spammer detection, misinformation, graph anomaly detection, deep learning, graph embedding, graph representation, graph neural networks

1 INTRODUCTION

ANOMALIES were first defined by Grubbs in 1969 [1] as “*one that appears to deviate markedly from other members of the sample in which it occurs*” and the studies on anomaly detection were initiated by the statistics community in the 19th century. To us, anomalies might appear as social spammers or

misinformation in social media; fraudsters, bot users or sexual predators in social networks; network intruders or malware in computer networks and broken devices or malfunctioning blocks in industry systems, and they often introduce huge damage to the real-world systems they appear in. According to FBI’s 2014 Internet Crime Report¹, the financial loss due to crime on social media reached more than \$60 million in the second half of the year alone and a more up-to-date report² indicates that the global economic cost of online fake news reached around \$78 billion a year in 2020.

In computer science, the research on anomaly detection dates back to the 1980s, and detecting anomalies on graph data has been an important data mining paradigm since the beginning. However, the extensive presence of connections between real-world objects and advances in graph data mining in the last decade have revolutionized our understanding of the graph anomaly detection problems such that this research field has received a dramatic increase in interest over the past five years. One of the most significant changes is that graph anomaly detection has evolved from relying heavily on human experts’ domain knowledge into machine

- Xiaoxiao Ma, Jia Wu, Shan Xue, Jian Yang, and Quan Z. Sheng are with the School of Computing, Macquarie University, Sydney, NSW 2109, Australia. E-mail: xiaoxiao.ma2@students, {jia.wu, emma.xue, jian.yang, michael.sheng}@mq.edu.au.
- Chuan Zhou is with the Academy of Mathematics and Systems Science, Chinese Academy of Sciences, Beijing 100190, China. E-mail: zhouchuan@amss.ac.cn.
- Hui Xiong is with the Department of Management Science and Information Systems, Rutgers, State University of New Jersey, Piscataway, NJ 08854 USA. E-mail: hxiong@rutgers.edu.
- Leman Akoglu is with the Heinz College of Information Systems and Public Policy, Carnegie Mellon University’s, Pittsburgh, PA 15213 USA. E-mail: lakoglu@andrew.cmu.edu.

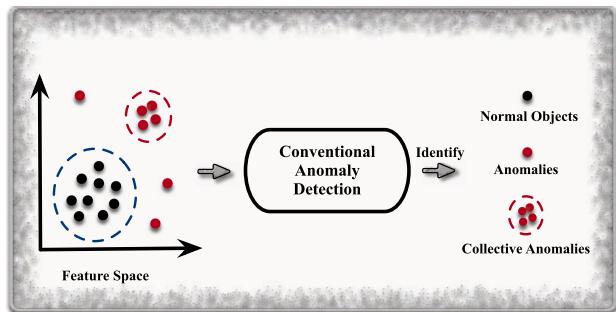
Manuscript received 19 April 2021; revised 20 August 2021; accepted 11 September 2021. Date of publication 8 October 2021; date of current version 8 November 2023. This work was supported by the Australian Research Council (ARC) DECRA Project No. DE200100964.

(Corresponding author: Jia Wu.)

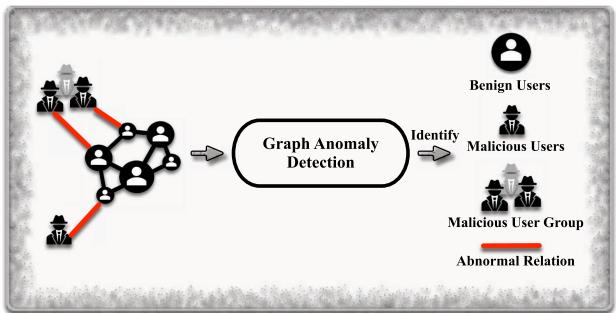
Recommended for acceptance by A. SI.

Digital Object Identifier no. 10.1109/TKDE.2021.3118815

1. https://www.fbi.gov/file-repository/2014_ic3report.pdf/view
2. <https://www.zdnet.com/article/online-fake-news-costing-us-78-billion-globally-each-year/>



(a) Conventional Anomaly Detection



(b) Graph Anomaly Detection

Fig. 1. Toy examples of conventional anomaly detection and graph anomaly detection. Apart from anomalies shown in (b), graph anomaly detection also identifies graph-level anomalies, detailed in Sections 8 and 9.

learning techniques that eliminate human intervention, and more recently, to various deep learning technologies. These deep learning techniques are not only capable of identifying potential anomalies in graphs far more accurately than ever before, but they can also do so in real-time.

For our purposes today, anomalies, which are also known as outliers, exceptions, peculiarities, rarities, novelties, etc., in different application fields, refer to abnormal objects that are significantly different from the standard, normal, or expected. Although these objects rarely occur in real-world, they contain critical information to support downstream applications. For example, the behaviors of fraudsters provide evidences for anti-fraud detection and abnormal network traffics reveal signals for network intrusion protection. Anomalies, in many cases, may also have real and adverse impacts, for instance, fake news in social media can create panic and chaos with misleading beliefs [2], [3], [4], [5], untrustworthy reviews in online review systems can affect customers' shopping choices [6], [7], [8], network intrusions might leak private personal information to hackers [9], [10], [11], [12], and financial frauds can cause huge damage to economic systems [13], [14], [15], [16].

Anomaly detection is the data mining process that aims to identify the unusual patterns that deviate from the majorities in a dataset [17], [18], [19]. In order to detect anomalies, conventional techniques typically represent real-world objects as feature vectors (e.g., news in social media are represented as bag-of-words [20], and images in web pages are represented as color histograms [21]), and then detect outlying data points in the vector space [22], [23], [24], as shown in Fig. 1(a). Although these techniques have shown power in locating deviating data points under tabulated data format, they inherently discard the complex relationships between objects [25].

Yet, in reality, many objects have rich relationships with each other, which can provide valuable complementary information for anomaly detection. Take online social networks as an example, fake users can be created using valid information from normal users or they can camouflage themselves by mimicking benign users' attributes [26], [27]. In such situations, fake users and benign users would have near-identical features, and conventional anomaly detection techniques might not be able to identify them using feature information only. Meanwhile, fake users always build relationships with a large number of benign users to increase their reputation and influence so they can get unexpected benefits, whereas benign users rarely exhibit such activities [28], [29]. Hence, these dense and unexpected connections formed by fake users denote their deviations to the benigns and more comprehensive detection techniques should take these structural information into account to pinpoint the deviating patterns of anomalies.

To represent the structural information, *Graphs*, in which nodes/vertices denote real objects, and the edges denote their relationships, have been prevalently used in a range of application fields [30], [31], [32], [33], [34], [35], including social activities, e-commerce, biology, academia and communication. With the structural information contained in graphs, detecting anomalies in graphs raises a more complex anomaly detection problem in non-Euclidean space - graph anomaly detection (GAD) that aims to identify anomalous graph objects (i.e., nodes, edges or sub-graphs) in a single graph as well as anomalous graphs among a set/database of graphs [25], [36], [37]. As a toy example shown in Fig. 1(b), given an online social network, graph anomaly detection aims to identify anomalous nodes (i.e., malicious users), anomalous edges (i.e., abnormal relations) and anomalous sub-graphs (i.e., malicious user groups). But, because the copious types of graph anomalies cannot be directly represented in euclidean feature space, it is not feasible to directly apply traditional anomaly detection techniques to graph anomaly detection, and researchers have intensified their efforts to GAD recently.

Amongst earlier works in this area, the detection methods relied heavily on handcrafted feature engineering or statistical models built by domain experts [38], [39], [40]. This inherently limits these techniques' capability to detect unknown anomalies, and the exercise tended to be very labor-intensive. Many machine learning techniques, such as matrix factorization [41], [42] and SVM [43], have also been applied to detect graph anomalies. However, real-world networks often contain millions of nodes and edges that result in extremely high dimensional and large-scale data, and these techniques do not easily scale up to such data efficiently. Practically, they exhibit high computational overhead in both the storage and execution time [44]. These general challenges associated with graph data are significant for the detection techniques, and we categorize them as data-specific challenges (Data-CHs) in this survey. A summary of them is provided in Appendix A, which can be found on the Computer Society Digital Library at <http://doi.ieeecomputersociety.org/10.1109/TKDE.2021.3118815>.

Non-deep learning based techniques also lack the capability to capture the non-linear properties of real objects [24]. Hence, the representations of objects learned by them are

not expressive enough to fully support graph anomaly detection. To tackle these problems, more recent studies seek the potential of adopting deep learning techniques to identify anomalous graph objects. As a powerful tool for data mining, deep learning has achieved great success in data representation and pattern recognition [55], [56], [57]. Its deep architecture with layers of parameters and transformations appear to suit the aforementioned problems well. The more recent studies, such as deep graph representation learning and graph neural networks (GNNs), further enrich the capability of deep learning for graph data mining [58], [59], [60], [61], [62]. By extracting expressive representations such that graph anomalies and normal objects can be easily separated, or the deviating patterns of anomalies can be learned directly through deep learning techniques, graph anomaly detection with deep learning (GADL) is starting to take the lead in the forefront of anomaly detection. As a frontier technology, graph anomaly detection with deep learning, hence, is expected to generate more fruitful results on detecting anomalies and secure a more convenient life for the society.

1.1 Challenges in GAD With Deep Learning

Due to the complexity of anomaly detection and graph data mining [63], [64], [65], [66], [67], in addition to the prior mentioned data-specific challenges, adopting deep learning techniques for graph anomaly detection also faces a number of challenges from the technical side. These challenges associated with deep learning are categorized as technique-specific challenges (Tech-CHs), and they are summarized as follows.

Tech-CH1. Anomaly-Aware Training Objectives. Deep learning models rely heavily on the training objectives to fine-tune all the trainable parameters. For graph anomaly detection, this necessitates appropriate training objectives or loss functions such that the GADL models can effectively capture the differences between benign and anomalous objects. Designing anomaly-aware objectives is very challenging because there is no prior knowledge about the ground-truth anomalies as well as their deviating patterns versus the majority. How to effectively separate anomalies from normal objects through training remains critical for deep learning-based models.

Tech-CH2. Anomaly Interpretability. In real-world scenarios, the interpretability of detected anomalies is also vital because we need to provide convincing evidence to support the subsequent anomaly handling process. For example, the risk management department of a financial organization must provide lawful evidence before blocking the accounts of identified anomalous users. As deep learning has been limited for its interpretability [24], [68], how to justify the detected graph anomalies remains a big challenge for deep learning techniques.

Tech-CH3. High Training Cost. Although D(G)NNs are capable of digesting rich information (e.g., structural information and attributes) in graph data for anomaly detection, these GADL models are more complex than conventional deep neural networks or machine learning methods due to the anomaly-aware training objectives. Such complexity inherently leads to high training costs in both time and computing resources.

Tech-CH4. Hyperparameter Tuning. D(G)NNs naturally exhibit a large set of hyperparameters, such as the number of neurons in each neural network layer, the learning rate, the weight decay and the number of training epochs. Their learning performance is significantly affected by the values of these hyperparameters. However, it remains a serious challenge to effectively select the optimal/sub-optimal settings for the detection models due to the lack of labeled data in real scenarios.

Because deep learning models are sensitive to their associated hyperparameters, setting well-performing values for the hyperparameters is vital to the success of a task. Tuning hyperparameter is relatively trivial in supervised learning when labeled data are available. For instance, users can find an optimal/sub-optimal set of hyperparameters (e.g., through random search, grid search) by comparing the model's outputs with the ground-truth. However, unsupervised anomaly detection has no accessible labeled data to judge the model's performance under different hyperparameter settings [69], [70]. Selecting the ideal hyperparameter values for unsupervised detection models persists as a critical obstacle to applying them in a wide range of real scenarios.

1.2 Existing Anomaly Detection Surveys

Recognizing the significance of anomaly detection, many review works have been conducted in the last ten years covering a range of anomaly detection topics: anomaly detection with deep learning, graph anomaly detection, graph anomaly detection with deep learning, and particular applications of graph anomaly detection such as social media, social networks, fraud detection and network security, etc.

There are some representative surveys on generalized anomaly detection techniques - [18], [45] and [44]. But only the most up-to-date work in Thudumu *et al.* [44] covers the topic of graph anomaly detection. Recognizing the power of deep learning, the three contemporary surveys, Ruff *et al.* [71], Pang *et al.* [24] and Chalapathy and Chawla [47] specifically review deep learning based anomaly detection techniques specifically.

As for graph anomaly detection, Akoglu *et al.* [25], Ranshous *et al.* [48], and Jennifer and Kumar [49] put their concentration on graph anomaly detection, reviewing many conventional approaches in this area, including statistical models and machine learning techniques. Other surveys are dedicated to particular applications of graph anomaly detection, such as computer network intrusion detection and anomaly detection in online social networks, e.g., [3], [6], [19], and [50], [51], [52], [53], [54]. These works provided solid reviews of the application of anomaly detection/graph anomaly detection techniques in these high demand and vital domains. However, none of the mentioned surveys are dedicated to techniques on graph anomaly detection with deep learning, as shown in Table 1, and hence do not provide a systematic and comprehensive review of these techniques.

1.3 Contributions

Our contributions are summarized as follows:

- *The first survey in graph anomaly detection with deep learning.* To the best of our knowledge, our survey is

TABLE 1
A Comparison Between Existing Surveys on Anomaly Detection

| Surveys | AD | DAD | GAD | GADL | | | Source Code | Dataset | |
|-------------------------------|----|-----|-----|------|------|-----------|-------------|------------|-----------|
| | | | | Node | Edge | Sub-graph | | Real-world | Synthetic |
| Our Survey | ● | ● | ● | ● | ● | ● | ● | ● | ● |
| Chandola <i>et al.</i> [18] | ● | - | - | - | - | - | - | - | - |
| Boukerche <i>et al.</i> [45] | ● | ● | - | - | - | - | - | - | - |
| Bulusu <i>et al.</i> [46] | ● | ● | - | - | - | - | - | - | - |
| Thudumu <i>et al.</i> [44] | ● | ● | ● | - | - | - | - | - | - |
| Pang <i>et al.</i> [24] | ● | ● | ● | ● | ● | - | ● | ● | - |
| Chalapathy and Chawla [47] | ● | ● | - | - | - | - | ● | ● | - |
| Akoglu <i>et al.</i> [25] | ● | - | ● | - | - | - | - | - | - |
| Ranshous <i>et al.</i> [48] | ● | - | ● | - | - | - | ● | - | - |
| Jennifer and Kumar [49] | ● | - | ● | - | - | - | - | - | - |
| Eltanbouly <i>et al.</i> [50] | ● | ● | ● | - | - | - | - | - | - |
| Fernandes <i>et al.</i> [51] | ● | ● | ● | - | - | - | - | - | - |
| Kwon <i>et al.</i> [52] | ● | ● | - | - | - | - | - | ● | - |
| Gogoi <i>et al.</i> [53] | ● | ● | - | - | - | - | - | - | - |
| Savage <i>et al.</i> [54] | ● | - | ● | - | - | - | - | - | - |
| Yu <i>et al.</i> [6] | ● | - | ● | - | - | - | - | - | - |
| Hunkelmann <i>et al.</i> [3] | ● | - | ● | - | - | - | - | - | - |
| Pourhabibi <i>et al.</i> [19] | ● | ● | ● | ● | - | - | - | - | - |

* AD: Anomaly Detection, DAD: Anomaly Detection with Deep Learning, GAD: Graph Anomaly Detection.

* GADL: Graph Anomaly Detection with Deep Learning.

* -: not included, ● (1-2 references included), ● (3-10 references included), ● (10+ references included).

We mark edge, sub-graph and graph detection as ● in our survey because we review more deep learning based works than any previous surveys.

the first to review the state-of-the-art deep learning techniques for graph anomaly detection. Most of the relevant surveys focus either on conventional graph anomaly detection methods using non-deep learning techniques or on generalized anomaly detection techniques (for tabular/point data, time series, etc.). Until now, there has been no dedicated and comprehensive survey on graph anomaly detection with deep learning. Our work bridges this gap, and we expect that an organized and systematic survey will help push forward research in this area.

- *A systematic and comprehensive review.* In this survey, we review the most up-to-date deep learning techniques for graph anomaly detection published in influential international conferences and journals in the area of deep learning, data mining, web services, and artificial intelligence, including: TKDE, TKDD, TPAMI, NeurIPS, SIGKDD, ICDM, WSDM, SDM, SIGMOD, IJCAI, AAAI, ICDE, CIKM, ICML, WWW, CVPR, and others. We first summarize seven data-specific and four technique-specific challenges in graph anomaly detection with deep learning. We then comprehensively review existing works from the perspectives of: 1) the motivations behind the deep methods; 2) the main ideas for identifying graph anomalies; 3) a brief introduction to conventional non-deep learning techniques; and 4) the technical details of deep learning algorithms. A brief timeline of graph anomaly detection and reviewed works is given in Fig. 2.
- *Future directions.* From the survey results, we highlight 12 future research directions covering emerging problems introduced by graph data, anomaly detection, deep learning models, and real-world applications. These future opportunities indicate challenges

that have not been adequately tackled, and so more effort is needed in the future.

- *Affluent resources.* Our survey also provides an extensive collection of open-sourced anomaly detection algorithms, public datasets, synthetic dataset generating techniques, as well as commonly used evaluation metrics to push forward the state-of-the-art in graph anomaly detection. These published resources offer benchmark datasets and baselines for future research.
- *A new taxonomy.* We have organized this survey with regard to different types of anomalies (i.e., nodes, edges, sub-graphs, and graphs) existing in graphs or graph databases. We also pinpoint the differences and similarities between different types of graph anomalies.

The rest of this survey is organized as follows. In Section 2, we provide preliminaries about the different types of settings. In Sections 3, 4, 5, 6, 7, 8, and 9, we review existing techniques for detecting anomalous nodes, edges, sub-graphs and graphs, respectively. In Section 10, we first provide a collection of published graph anomaly detection algorithms and datasets and then summarize commonly used evaluation metrics and synthetic data generation strategies. We highlight 12 future directions concerning deep learning in graph anomaly detection in Section 11 and summarize our survey in Section 12. A concrete taxonomy of our survey is given in Appendix B, available in the online supplemental material.

2 PRELIMINARIES

In this section, we provide definitions of different types of graphs mostly used in node/edge/sub-graph-level anomaly detection (Sections 3, 4, 5, 6, and 7). For

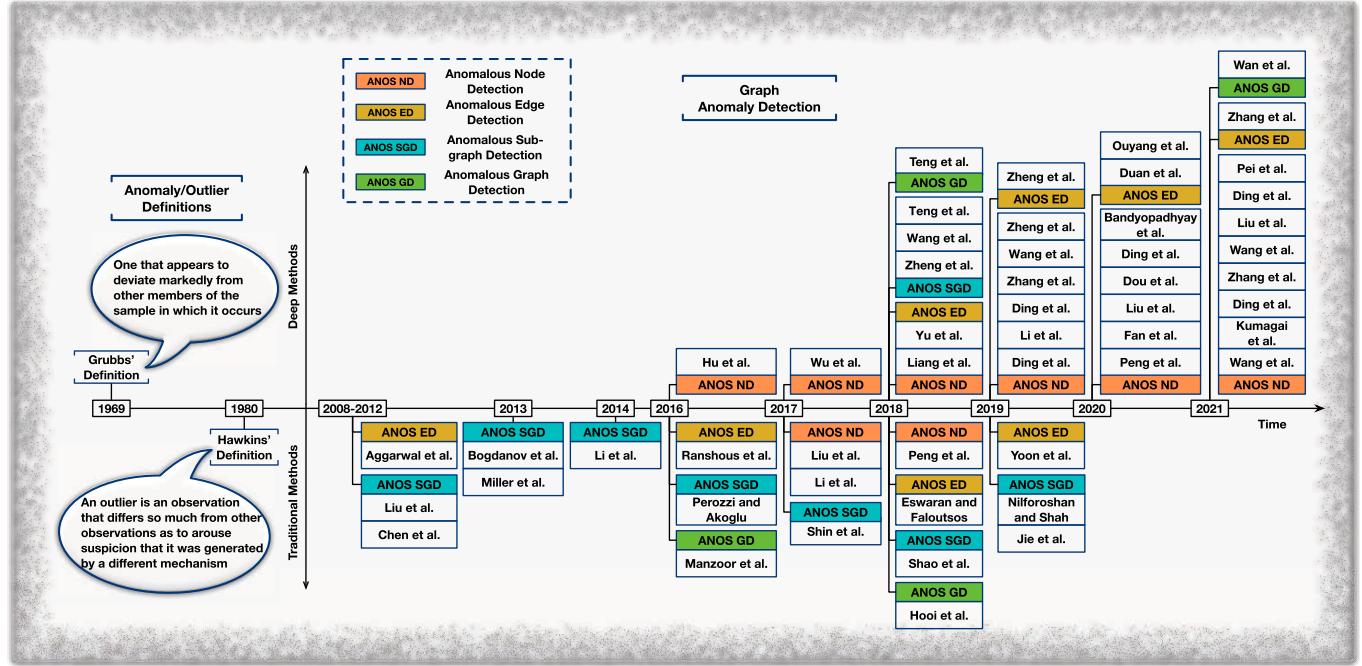


Fig. 2. A timeline of graph anomaly detection and reviewed techniques.

consistency, we have followed the conventional categorization of graphs as in existing works [25], [48], [52] and categorize them as static graphs, dynamic graphs, and graph databases. Unless otherwise specified, all graphs mentioned in the following sections are static. Meanwhile, as graph-level anomaly detection is discussed far away on page 13, to enhance readability, the definition for the graph database is given closer to the material in Section 8.

Definition 1. (Plain Graph). A static plain graph $G = \{V, E\}$ comprises a node set $V = \{v_i\}_1^n$ and an edge set $E = \{e_{i,j}\}$ where n is the number of nodes and $e_{i,j} = (v_i, v_j)$ denotes an edge between nodes v_i and v_j . The adjacency matrix $A = [a_{i,j}]_{n \times n}$ restores the graph structure, where $a_{i,j} = 1$ if node v_i and v_j is connected, otherwise $a_{i,j} = 0$.

Definition 2. (Attributed Graph). A static attributed graph $G = \{V, E, X\}$ comprises of a node set V , an edge set E and an attribute set X . In an attributed graph, the graph structure follows the definition in Definition 1. The attribute matrix $X = [\mathbf{x}_i]_{n \times k}$ consists of nodes' attribute vectors, where \mathbf{x}_i is the attribute vector associated with node v_i and k is the vector's dimension. Hereafter, the terms attribute and feature are used interchangeably.

Definition 3. (Dynamic Graph). A dynamic graph $G(t) = \{V(t), E(t), X_v(t), X_e(t)\}$ comprises nodes and edges changing overtime. $V(t)$ is the node set in the graph at a specific time step t , $E(t)$ is the corresponding edge set, $X_v(t)$ and $X_e(t)$ are the node attribute matrix and edge attribute matrix at time step t in the graph if existed.

In reality, the nodes or edges might also be associated with numerical or categorical labels to indicate their classes (e.g., normal or abnormal). When label information is available/partially-available, supervised/semi-supervised detection models could be effectively trained.

3 ANOMALOUS NODE DETECTION (ANOS ND)

Anomalous nodes are commonly recognized as individual nodes that are significantly different from others. In real-world applications, these nodes often represent abnormal objects that appear individually, such as a single network intruder in computer networks, an independent fraudulent user in online social networks or a specific fake news on social media. In this section, we specifically focus on anomalous node detection in static graphs. The reviews on dynamic graphs can be found in Section 4. Table 2 provides a summary of techniques reviewed for ANOS ND.

When detecting anomalous nodes in static graphs, the differences between anomalies and regular nodes are mainly drawn from the graph structural information and nodes/edges' attributes [41], [72], [73], [74]. Given prior knowledge (i.e., community structure, attributes) about a static graph, anomalous nodes can be further categorized into the following three types:

- *Global anomalies* only consider the node attributes. They are nodes that have attributes significantly different from all other nodes in the graph.
- *Structural anomalies* only consider the graph structural information. They are abnormal nodes that have different connection patterns (e.g., connecting different communities, forming dense links with others).
- *Community anomalies* consider both node attributes and graph structural information. They are defined as nodes that have different attribute values compared to other nodes in the same community.

In Fig. 3, node 14 is a global anomaly because its 4th feature value is 1 while all other nodes in the graph have the value of 0 for the corresponding feature. Nodes 5, 6, and 11 are identified as structural anomalies because they have links with other communities while other nodes in their

TABLE 2
Summary of Anomalous Node Detection Techniques

| Graph Type | Approach | Category | Objective Function | Measurement | Outputs |
|----------------------------|------------------|----------|---|-----------------------------|---|
| Static Graph - Plain | [77] | NR | $\sum_{(i,j) \in E} \ \mathbf{Z}_i - \mathbf{Z}_j\ ^2 + \alpha \sum_{(i,j) \notin E} (\ \mathbf{Z}_i - \mathbf{Z}_j\ - 1)^2$ | Anomaly Score | $\sum_{k=1}^d \frac{y_i^k}{y_i}$ |
| | DCI [105] | NR | $\frac{1}{K} \sum_{k=1}^K \mathcal{L}_{DCI}^k$ | Anomaly Prediction | Predicted Label |
| | NAC [102] | RL | Cumulative reward | - | Anomalies |
| Static Graph - Attributed | ALAD [106] | Non-DP | $\min_{W,H} \ A - WW^T\ _F^2 + \alpha \ X - WH\ _F^2 + \gamma (\ W\ _F^2 + \ H\ _F^2)$ | Anomaly Score | $\frac{\mathbf{W}_{n,c}}{\sum_c \mathbf{W}_{n,c}} \cos(\mathbf{A}_{n*}, \mathbf{H}_{c*})$ |
| | Radar [41] | Non-DP | $\min_{W,R} \ X - W^T X - R\ _F^2 + \alpha \ W\ _{2,1} + \beta \ R\ _{2,1} + \gamma \text{tr}(R^T L R)$ | Residual Analysis | Residual Value |
| | ANOMALOUS [107] | Non-DP | $\min_{W,\tilde{R}} \ X - XWX - \tilde{R}\ _F^2 + \alpha \ W\ _{2,1} + \beta \ W^T\ _{2,1} + \gamma \ \tilde{R}^T\ _{2,1} + \phi \text{tr}(\tilde{R}LR^T)$ | Residual Analysis | Residual Value |
| | SGASD [108] | Non-DP | $\min_{\mathbf{w}, \mathbf{c}} \frac{1}{2} \sum_{i=1}^m \mathbf{c}_i (V_{i,*} \mathbf{w} - y_i)^2 + \frac{\lambda_1}{2} \ \mathbf{w}\ _2^2 + \lambda_2 \sum_{j=1}^d \ \mathbf{c}_{G_j^i}\ _2$ | Anomaly Prediction | Predicted Label |
| | DONE [82] | DNN | $\frac{\alpha_1 \mathcal{L}_{str}^{Recs} + \alpha_2 \mathcal{L}_{attr}^{Recs} + \alpha_3 \mathcal{L}_{str}^{Hom} + \alpha_4 \mathcal{L}_{attr}^{Hom} + \alpha_5 \mathcal{L}_{Com}^{Com}}{\alpha_1 \mathcal{L}_{str}^{Hom} + \alpha_2 \mathcal{L}_{attr}^{Hom}}$ | Anomaly Scores | o_i^s, o_i^a, o_i^{com} |
| | DOMINANT [92] | GCN | $(1 - \alpha) \mathcal{R}_S + \alpha \mathcal{R}_A$ | Anomaly Score | $(1 - \alpha) \ \mathbf{a}_i - \hat{\mathbf{a}}_i\ _2 + \alpha \ \mathbf{x}_i - \hat{\mathbf{x}}_i\ _2$ |
| | ALARM [93] | GCN | $\sum_{i=1}^n \sum_{j=1}^n -[\gamma A_{ij} \log \hat{A}_{ij} + (1 - A_{ij}) \log(1 - \hat{A}_{ij})] + \mathcal{L}_a$ | Anomaly Score | $(1 - \alpha) \ \mathbf{a}_i - \hat{\mathbf{a}}_i\ _2^2 + \alpha \ \mathbf{x}_i - \hat{\mathbf{x}}_i\ _2^2$ |
| | SpecAE [97] | GCN | $\mathbb{E}[dis(X, \hat{X})] + \mathbb{E}[dis(X, \tilde{X})] + \lambda_1 \mathbb{E}(E(Z)) + \lambda_2 KL$ | Density Estimation | Anomalousness Rank |
| Dynamic Graph - Plain | Fdgars [98] | GCN | \mathcal{L}_{GCN} | Anomaly Prediction | Predicted Label |
| | GraphRfi [99] | GCN | $\mathcal{L}_{rating} + \lambda \mathcal{L}_{fraudster}$ | Anomaly Prediction | Predicted Label |
| | ResGCN [109] | GCN | $(1 - \alpha) \ A - \hat{A}\ _F^2 + \alpha \ X - \hat{X} - \lambda R\ _F^2$ | Anomaly Score | $\ R_{i,:}\ _2$ |
| | GraphUCB [100] | RL | Expert Judgment | - | Anomalies |
| | AnomalyDAE [110] | GAT | $\alpha \ (A - \hat{A}) \odot \theta\ _F^2 + (1 - \alpha) \ (X - \hat{X}) \odot \eta\ _F^2$ | Reconstruction Loss | Anomalousness Rank |
| | SemiGNN [111] | GAT | $\alpha \mathcal{L}_{sup} + (1 - \alpha) \mathcal{L}_{unsup} + \lambda \mathcal{L}_{reg}$ | Anomaly Prediction | Predicted Label |
| | AEGIS [112] | GAN | $\mathcal{L}_{AE} + \mathcal{L}_{GAN}$ | Anomaly Score | $1 - D(\mathbf{z}_i)$ |
| | REMAD [113] | NR | $\mathcal{L}_{res} + \beta \ R^T\ _{2,1}$ | Residual Analysis | Residual Value |
| | CARE-GNN [27] | NR | $\mathcal{L}_{GNN} + \lambda_1 \mathcal{L}_{Simi}^{(1)} + \lambda_2 \mathcal{L}_{reg}$ | Anomaly Prediction | Predicted Label |
| | SEANO [114] | NR | $-\sum_{i \in V_L} \log p(y_i \mathbf{x}_i, \bar{\mathbf{x}}_{N_i}) - \sum_{i \in V} \sum_{v' \in C_i} \log p(v' \mathbf{x}_i, \bar{\mathbf{x}}_{N_i})$ | Anomaly Score | Discriminator's Output |
| | OCGNN [115] | NR | $\frac{1}{\beta K} \sum_{v_i \in \mathbf{V}_{tr}} [\ g(X, A; \mathcal{W})_{v_i} - c\ ^2 - r^2] + r^2 + \frac{\lambda}{2} \sum_{l=1}^L \ W^{(l)}\ ^2$ | Location in Embedding Space | Distance to Hypersphere Center |
| | GAL [67] | NR | $\max\{0, \max_{y_{v'} \neq y_u} g(u, v') - \min_{y_v = y_u} g(u, v) + \Delta_{y_u}\}$ | Anomaly Prediction | Predicted Label |
| | CoLA [116] | NR | $-\sum_{i=1}^N y_i \log(CLM(v_i, \mathcal{G}_i)) + (1 - y_i) \log(1 - CLM(v_i, \mathcal{G}_i))$ | Anomaly Score | $\frac{\sum_{r=1}^R (s_{i,r}^{(-)} - s_{i,r}^{(+)})}{R}$ |
| | COMMANDER [117] | NR | $-\mathcal{L}_D + \mathcal{L}_C + \mathcal{L}_R$ | Anomaly Score | $\bar{y}_i \ \bar{\mathbf{x}}_i - \mathbf{x}_i\ _2^2$ |
| | FRAUDRE [118] | NR | $\sum_{i=1}^n f^*(y_i, \mathbf{h}_i^{(final)} \mathbf{W}_2)$ | Anomaly Prediction | Predicted Label |
| | Meta-GDN [119] | NR | $(1 - y_i) \cdot dev(v_i) + y_i \cdot \max(0, dev(v_i))$ | Anomaly Score | $\mathbf{u}_s^T \mathbf{o}_i + b_s$ |
| Dynamic Graph - Plain | NetWalk [84] | DNN | $\gamma \mathcal{L}_{AE} + \mathcal{L}_{Clique} + \lambda \ W\ _F^2 + \beta KL$ | Anomaly Score | Nearest Distance to Cluster Centers |
| Dynamic Graph - Attributed | MTHL [120] | Non-DP | $\min_{\mathcal{P}} f(\mathcal{P})$ | Anomaly Score | Distance to Hypersphere Centroid |
| | OCAN [121] | GAN | $\mathcal{L}_{LSTM-AE} + \mathcal{L}_{GAN}$ | Anomaly Score | Discriminator's Output |

* Non-DP: Non-Deep Learning Techniques, DNN: Deep NN Based Techniques, GCN: GCN Based Techniques, RL: Reinforcement Learning Based Techniques.
* GAT: GAT Based Techniques, NR: Network Representation Based Techniques, GAN: Generative Adversarial Network Based Techniques.

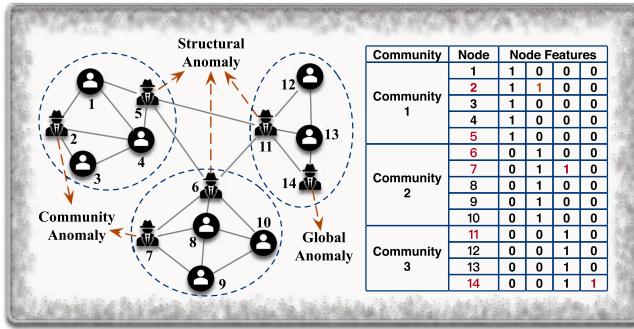


Fig. 3. Three types of anomalous nodes: Structural anomalies, community anomalies and global anomalies.

communities do not form cross-community links. Nodes 2 and 7 are community anomalies because their feature values are different from others in the communities they belong to.

3.1 ANOS ND on Plain Graphs

Plain graphs are dedicated to representing the structural information in real-world networks. To detect anomalous nodes in plain graphs, the graph structure has been extensively exploited from various angles. Here, we first summarize the representative traditional non-deep learning approaches, followed by more recent, advanced detection techniques based on representation learning and reinforcement learning.

3.1.1 Traditional Non-Deep Learning Techniques

Prior to the recent advances in deep learning and other state-of-the-art data mining technologies, traditional non-deep learning techniques have been widely used in many real-world networks to identify anomalous entities. A key idea behind these techniques was to transform the graph anomaly detection into a traditional anomaly detection problem, because the graph data with rich structure information can not be handled by the traditional detection techniques (for tabular data only) directly. To bridge the gap, many approaches [38], [75], [76] used the statistical features associated with each node, such as in/out degree, to detect anomalous nodes.

For instance, OddBall [38] employs the statistical features (e.g., the number of 1-hop neighbors and edges, the total weight of edges) extracted from each node and its 1-hop neighbors to detect particular structural anomalies that: 1) form local structures in shape of near-cliques or stars; 2) have heavy links with neighbors such that the total weight is extremely large; or 3) have a single dominant heavy link with one of the neighbors.

With properly selected statistical features, anomalous nodes can be identified with respect to their deviating feature patterns. But, in real scenarios, it is very hard to choose the most suitable features from a large number of candidates, and domain experts can always design new statistics, e.g., the maximum/minimum weight of edges. As a result, these techniques often carry prohibitive cost for assessing the most significant features and do not effectively capture the structural information.

3.1.2 Network Representation Based Techniques

To capture more valuable information from the graph structure for anomaly detection, network representation techniques have been widely exploited. Typically, these techniques encode the graph structure into an embedded vector space and identify anomalous nodes through further analysis. Hu *et al.* [77], for example, proposed an effective embedding method to detect structural anomalies that are connecting with many communities. It first adopts a graph partitioning algorithm (e.g., METIS [78]) to group nodes into d communities (d is a user-specified number). Then, the method employs a specially designed embedding procedure to learn node embeddings that could capture the link information between each node and d communities. Denoting the embedding for node i as $Z_i = \{z_i^1, \dots, z_i^d\}$, the procedure initializes each $z_i^c \in Z_i$ with regard to the membership of node i to community c (if node i belongs to the community, then $z_i^c = \frac{1}{\sqrt{2}}$; otherwise, 0.) and optimizes node embeddings such that directly linked nodes have similar embeddings and unconnected nodes are dissimilar.

After generating the node embeddings, the link information between node i and d communities is quantified for further anomaly detection analysis. For a given node i , such information is represented as:

$$\overline{NB(i)} = (y_i^1, \dots, y_i^d) = \sum_{j \in NB(i)} (1 - \|Z_i - Z_j\|) \cdot Z_j, \quad (1)$$

where $NB(i)$ comprises node i 's neighbors. If i has many links with community c , then the value in the corresponding dimension y_i^c will be large.

In the last step, Hu *et al.* [77] formulated a scoring function to assign anomalousness scores, calculated as:

$$AScore(i) = \sum_{k=1}^d \frac{y_i^k}{y_i^*}, y_i^* = \max\{y_i^1, \dots, y_i^d\}. \quad (2)$$

As expected, structural anomalies receive higher scores as they connect to different communities. Indeed, given a predefined threshold, nodes with above-threshold scores are identified as anomalies.

To date, many plain network representation methods such as Deepwalk [79], Node2Vec [80] and LINE [81] have shown their effectiveness in generating node representations and been used for anomaly detection performance validation [82], [83], [84], [85]. By pairing the conventional anomaly detection techniques such as density-based techniques [86] and distance-based techniques [87] with node embedding techniques, anomalous nodes can be identified with regard to their distinguishable locations (i.e., low-density areas or far away from the majorities) in the embedding space.

3.1.3 Reinforcement Learning Based Techniques

The success of reinforcement learning (RL) in tackling real-world decision making problems has attracted substantial interests from the anomaly detection community because detecting anomalous nodes can be naturally regarded as a problem of deciding which class a node belongs to - anomalous or benign. In a recent work [102], Morales *et al.*

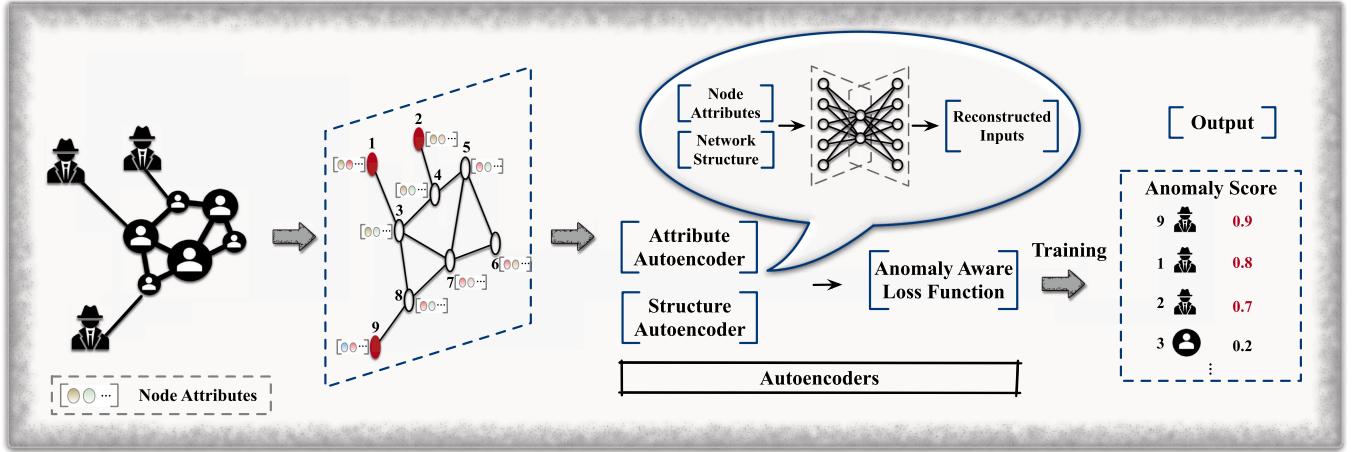


Fig. 4. ANOS ND on attributed graphs – Deep NN based approaches. As an example, the autoencoder is used to capture the graph structure and node attributes. With a specially-designed anomaly aware loss function, anomaly scores will be assigned to every node, and the top- k nodes are anomalies (e.g., nodes 9, 1, and 2 at top-3).

intuitively combined reinforcement learning and network embedding techniques for anomalous node detection. The proposed model, NAC, is trained with labeled data without any human intervention. Specifically, it first selects a seed network consisting of partially observed nodes and edges. Then, starting from the seed network, NAC adopts reinforcement learning to learn a node selection plan such that anomalous nodes in the undiscovered area can be identified. This is achieved by rewarding selection plans that can choose labeled anomalies with higher gains. Through offline training, NAC will learn an effective anomalous node selection strategy and discover potential anomalies in the undiscovered graph step by step.

3.2 ANOS ND on Attributed Graphs

In addition to the structural information, real-world networks also contain rich attribute information affiliated with nodes [88], [89]. These attributes provide complementary information about real objects and together with graph structure, more hidden anomalies that are non-trivial can now be detected.

For clarity, we distinguish between deep neural networks and graph neural networks in this survey. We review deep neural network (Deep NN) based techniques, GCN based techniques, and reinforcement learning based techniques for ANOS ND as follows. Due to page limitations, other existing works including traditional non-deep learning techniques, GAT [90] based techniques, GAN based techniques, and network representation based techniques are surveyed in Appendix C, available in the online supplemental material.

3.2.1 Deep NN Based Techniques

The deep learning models such as autoencoder and deep neural networks provide solid basis for learning data representations. Adopting these models for more effective anomalous node detection have drawn substantial interest recently.

For example, Bandyopadhyay *et al.* [82] developed an unsupervised deep model, DONE, to detect global anomalies, structural anomalies and community anomalies in attributed graphs. Specifically, this work measures three

anomaly scores for each node that indicate the likelihood of the situations where 1) it has similar attributes with nodes in different communities (o_i^a); or 2) it connects with other communities (o_i^s); or 3) it belongs to one community structurally but the attributes follow the pattern of another community (o_i^{com}). If a particular node exhibits any of these characteristics, then it is assigned a higher score and is anomalous.

To acquire these scores, DONE adopts two separate autoencoders (AE), i.e., a structure AE and an attribute AE, as shown in Fig. 4. Both are trained by minimizing the reconstruction errors and preserving the homophily that assumes connected nodes have similar representations in the graph. When training the AEs, nodes exhibiting the pre-defined characteristics are hard to reconstruct and therefore introduce more reconstruction errors because their structure or attribute patterns do not conform to the standard behavior. Hence, the adverse impact of anomalies should be alleviated to achieve the minimized error. Accordingly, DONE specially designs an anomaly-aware loss function with five terms: \mathcal{L}_{str}^{Recs} , $\mathcal{L}_{attr}^{Recs}$, \mathcal{L}_{str}^{Hom} , \mathcal{L}_{attr}^{Hom} , and \mathcal{L}_{com}^{Hom} . \mathcal{L}_{str}^{Recs} and $\mathcal{L}_{attr}^{Recs}$ are the structure reconstruction error and attribute reconstruction error that can be written as:

$$\mathcal{L}_{str}^{Recs} = \frac{1}{N} \sum_{i=1}^N \log \left(\frac{1}{o_i^s} \right) \|\mathbf{t}_i - \hat{\mathbf{t}}_i\|_2^2, \quad (3)$$

and

$$\mathcal{L}_{attr}^{Recs} \frac{1}{N} \sum_{i=1}^N \log \left(\frac{1}{o_i^a} \right) \|\mathbf{x}_i - \hat{\mathbf{x}}_i\|_2^2, \quad (4)$$

where N is the number of nodes, \mathbf{t}_i and \mathbf{x}_i store the structure information and attributes of node i , $\hat{\mathbf{t}}_i$ and $\hat{\mathbf{x}}_i$ are the reconstructed vectors. \mathcal{L}_{str}^{Hom} and \mathcal{L}_{attr}^{Hom} are proposed to maintain the homophily and they are formulated as:

$$\mathcal{L}_{str}^{Hom} = \frac{1}{N} \sum_{i=1}^N \log \left(\frac{1}{o_i^s} \right) \frac{1}{|N(i)|} \sum_{j \in N(i)} \|\mathbf{h}_i^s - \mathbf{h}_j^s\|_2^2, \quad (5)$$

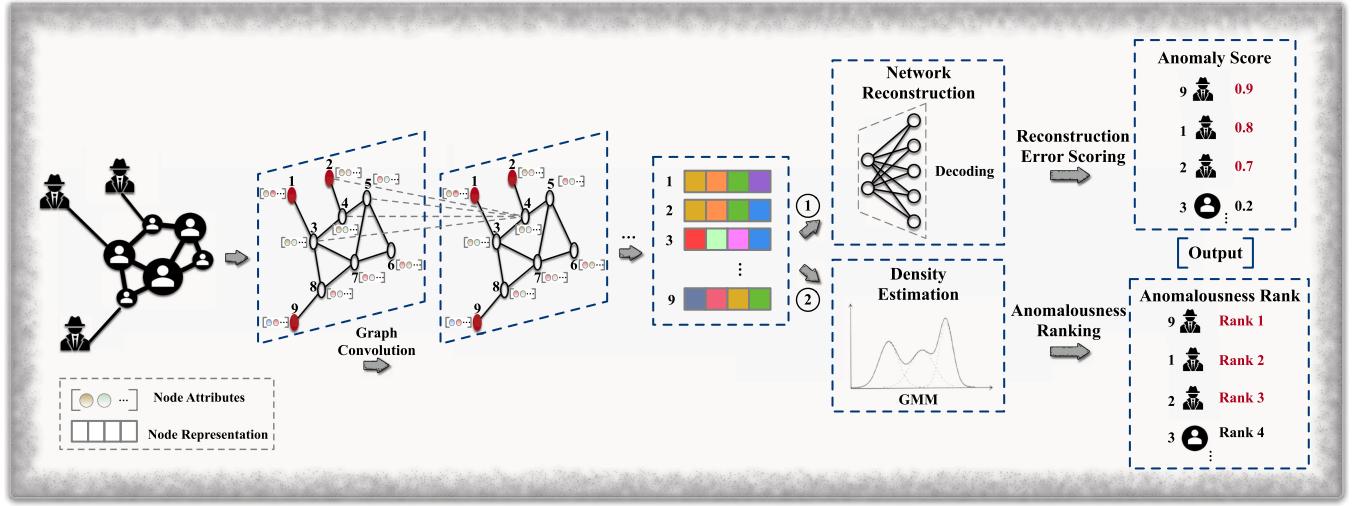


Fig. 5. ANOS ND on attributed graphs – GCN based approaches. Node representations are generated through GCN layers. Anomalies are then detected according to their reconstruction loss (①) or embedding distribution in the embedding space (②).

and

$$\mathcal{L}_{attr}^{Hom} = \frac{1}{N} \sum_{i=1}^N \log \left(\frac{1}{o_i^a} \right) \frac{1}{|N(i)|} \sum_{j \in N(i)} \|\mathbf{h}_i^a - \mathbf{h}_j^a\|_2^2, \quad (6)$$

where \mathbf{h}_i^s and \mathbf{h}_i^a are the learned latent representations from the structure AE and attribute AE, respectively. \mathcal{L}^{Com} poses further restrictions on the generated representations for each node by the two AEs such that the graph structure and node attributes complement each other. It is formulated as:

$$\mathcal{L}^{Com} = \frac{1}{N} \sum_{i=1}^N \log \left(\frac{1}{o_i^{com}} \right) \|\mathbf{h}_i^s - \mathbf{h}_i^a\|_2^2, \quad (7)$$

By minimizing the sum of these loss functions, the anomaly scores of each node are quantified, and the top- k nodes with higher scores are identified as anomalies.

3.2.2 GCN Based Techniques

Graph convolutional neural networks (GCNs) [91] have accomplished decent success in many graph data mining tasks (e.g., link prediction, node classification, and recommendation) owing to its capability of capturing comprehensive information in the graph structure and node attributes. Therefore, many anomalous node detection techniques start to investigate GCNs. Fig. 5 illustrates a general framework of existing works in this line.

In [92], Ding *et al.* measured an anomaly score for each node using the network reconstruction errors of both the structure and attribute. The proposed method, DOMINANT, comprises three parts, namely, the graph convolutional encoder, the structure reconstruction decoder, and the attribute reconstruction decoder. The graph convolutional encoder generates node embeddings through multiple graph convolutional layers. The structure reconstruction decoder tends to reconstruct the network structure from the learned node embeddings, while the attribute reconstruction decoder reconstructs the node attribute matrix. The whole neural network is trained to minimize the following loss function:

$$\begin{aligned} \mathcal{L}_{DOMINANT} &= (1 - \alpha)\mathcal{R}_S + \alpha\mathcal{R}_A \\ &= (1 - \alpha)\|A - \hat{A}\|_F^2 + \alpha\|X - \hat{X}\|_F^2, \end{aligned} \quad (8)$$

where α is the coefficient, A depicts the adjacency matrix of the graph, \mathcal{R}_S and \mathcal{R}_A quantify reconstruction errors with regard to the graph structure and node attributes, respectively. When the training is finished, an anomaly score is then assigned to each node according to its contribution to the total reconstruction error, which is calculated by:

$$score(i) = (1 - \alpha)\|\mathbf{a}_i - \hat{\mathbf{a}}_i\|_2 + \alpha\|\mathbf{x}_i - \hat{\mathbf{x}}_i\|_2, \quad (9)$$

where \mathbf{a}_i and \mathbf{x}_i are the structure vector and attribute vector of node i , $\hat{\mathbf{a}}_i$ and $\hat{\mathbf{x}}_i$ are their corresponding reconstructed vectors. The nodes are then ranked according to their anomaly scores in descending order, and the top- k nodes are recognized as anomalies.

To enhance the performance of anomalous node detection, later work by Peng *et al.* [93] further explores node attributes from multiple attributed views to detect anomalies. The multiple attributed views are employed to describe different perspectives of the objects [94], [95], [96]. For example, in online social networks, user's demographic information and posted contents are two different attributed views, and they characterize the personal information and social activities, respectively. The underlying intuition of investigating different views is that anomalies might appear to be normal in one view but abnormal in another view.

For the purpose of capturing these signals, the proposed method, ALARM, applies multiple GCNs to encode information in different views and adopts a weighted aggregation of them to generate node representations. This model's training strategy is similar to DOMINANT [92] in that it aims to minimize the network reconstruction loss and attribute reconstruction loss and can be formulated as:

$$\begin{aligned} \mathcal{L}_{ALARM} &= \sum_{i=1}^n \sum_{j=1}^n -[\gamma A_{ij} \log \hat{A}_{ij} + (1 - A_{ij}) \log (1 - \hat{A}_{ij})] \\ &\quad + \|X - \hat{X}\|_2^F, \end{aligned} \quad (10)$$

where γ is coefficient to balance the errors, A_{ij} is the element at coordinate (i, j) in the adjacency matrix A , \hat{A}_{ij} is the corresponding element in the reconstructed adjacency matrix \hat{A} , X is the original node feature matrix and \hat{X} is the reconstructed node feature matrix. Lastly, ALARM adopts a similar scoring function as [92], and nodes with top- k highest scores are anomalous.

Instead of spotting unexpected nodes using their reconstruction errors, Li *et al.* [97] proposed SpecAE to detect global anomalies and community anomalies via a density estimation approach, Gaussian Mixture Model (GMM). Global anomalies can be identified by only considering the node attributes. For community anomalies, the structure and attributes need to be jointly considered because of their distinctive attributes to the neighbors. Accordingly, SpecAE investigates a graph convolutional encoder to learn node representations and reconstruct the nodal attributes through a deconvolution decoder. The parameters in the GMM are then estimated using the node representations. Due to the deviating attribute patterns of global and community anomalies, normal nodes are expected to exhibit greater energies in GMM, and the k nodes with the lowest probabilities are deemed to be anomalies.

In [98], Wang *et al.* developed a novel detection model that identify fraudsters using their relations and features. Their proposed method, Fdgars, first models online users' reviews and visited items as their features, and then identifies a small portion of significant fraudsters based on these features. In the last step, a GCN is trained in a semi-supervised manner by using the user-user network, user features, and labeled users. After training, the model can directly label unseen users.

A more recent work, GraphRfi [99], also explores the potential of combining anomaly detection with other downstream graph analysis tasks. It targets on leveraging anomaly detection to identify malicious users and provide more accurate recommendations to service benign users by alleviating the impact of these untrustworthy users. Specifically, a GCN framework is deployed to encode users and items into a shared embedding space for recommendation and users are classified as fraudsters or normal users through an additional neural random forest using their embeddings. For rating prediction between users and items, the framework reduces the corresponding impact of suspicious users by assigning less weights to their training loss. At the same time, the rating behavior of users also provides auxiliary information for fraudster detection. The mutually beneficial relationship between these two applications (anomaly detection and recommendation) indicates the potential of information sharing among multiple graph learning tasks.

3.2.3 Reinforcement Learning Based Techniques

In contrast to NAC [102] reviewed in Section 3.1.3, Ding *et al.* [100] investigated to the use of reinforcement learning for anomalous node detection in attributed graphs. Their proposed algorithm, GraphUCB, models both attribute information and structural information, and inherits the merits of the contextual multi-armed bandit technology [101] to output potential anomalies. By grouping nodes into k clusters based on their features, GraphUCB forms a k -armed

bandit model and measures the payoff of selecting a specific node as a potential anomaly for expert evaluation. With experts' feedback on the predicted anomalies, the decision-making strategy is continuously optimized. Eventually, the most potential anomalies can be selected.

4 ANOS ND ON DYNAMIC GRAPHS

Real-world networks can be modeled as dynamic graphs to represent evolving objects and the relationships among them. In addition to structural information and node attributes, dynamic graphs also contain rich temporal signals [103], e.g., the evolving patterns of the graph structure and node attributes. On the one hand, these information inherently makes anomalous node detection on dynamic graphs more challenging. This is because dynamic graphs usually introduce large volume of data and temporal signals should also be captured for anomaly detection. But, on the other hand, they could provide more details about anomalies [25], [48], [104]. In fact, some anomalies might appear to be normal in the graph snapshot at each time stamp, and, only when the changes in a graph's structure are considered, do they become noticeable.

In this section, we review the network representation based techniques and GAN based techniques as follows. Relevant techniques from traditional non-deep learning approaches are reviewed in Appendix D, available in the online supplemental material.

4.1 Network Representation Based Techniques

Following the research line of encoding a graph into an embedding space, after which anomaly detection is performed, dynamic network representation techniques have been investigated in the more recent works. Specifically, in [84], Yu *et al.* presented a flexible deep representation technique, called NetWalk, for detecting anomalous nodes in dynamic (plain) graphs using only the structure information. It adopts an autoencoder to learn node representations on the initial graph and incrementally updates them when new edges are added or existing edges are deleted. To detect anomalies, NetWalk first executes the streaming k -means clustering algorithm [122] to group existing nodes in the current time stamp into different clusters. Then, each node's anomaly score is measured with regard to its closest distance to the k clusters. When the node representations are updated, the cluster centers and anomaly scores are recalculated accordingly.

4.2 GAN Based Techniques

In practice, anomaly detection is facing great challenges from the shortage of ground-truth anomalies. Consequently, many research efforts have been invested in modeling the features of anomalies or regular objects such that anomalies can be identified effectively. Among these techniques, generative adversarial networks (GAN) [123] have received extensive attention because of its impressive performance in capturing real data distribution and generating simulated data.

Motivated by the recent advances in "bad" GAN [124], Zheng *et al.* [121] circumvented the fraudster detection problem using only the observed benign users' attributes. The basic

TABLE 3
Summary of Anomalous Edge, Sub-Graph and Graph Detection Techniques

| Graph Type | Approach | Category | Objective Function | Measurement | Outputs |
|--|------------------|----------|--|-------------------------------|--|
| Anomalous Edge Detection Techniques | | | | | |
| Static Graph - Plain | UGED [127] | DNN | cross-entropy($f(u, N(u)), v)$ | Anomaly Score | mean(1 – $P(v u, N(u)), 1 – P(u v, N(v))$) |
| | AANE [128] | GCN | $\mathcal{L} = \mathcal{L}_{afl} + \gamma \mathcal{L}_{aal}$ | Anomaly Ranking | Edge Existing Probability |
| Static Graph - Attributed | eFraudCom [146] | NR | $\mathcal{L}_{MAIN} - \lambda \mathcal{L}_{MI}$ | Anomaly Prediction | Predicted Label |
| Dynamic Graph - Plain | NetWalk [84] | NR | $\gamma \mathcal{L}_{AE} + \mathcal{L}_{Clique} + \lambda \ W\ _F^2 + \beta KL$ | Anomaly Score | Nearest Distance to Cluster Centers |
| Dynamic Graph - Attributed | AddGraph [131] | GCN | $\min \sum_{e \in \varepsilon} \sum_{e' \notin \varepsilon} \max\{0, \gamma + f(i, j, w) - f(i', j', w)\} + \lambda \mathcal{L}_{reg}$ | Anomaly Score | $f(i, j, w) = w \cdot \sigma(\beta \cdot (\ \mathbf{a} \odot \mathbf{h}_i + \mathbf{b} \odot \mathbf{h}_j\ _2^2 - \mu))$ |
| Anomalous Sub-graph Detection Techniques | | | | | |
| Static Graph - Plain | DeepFD [134] | NR | $\mathcal{L}_{recon} + \alpha \mathcal{L}_{sim} + \gamma \mathcal{L}_{reg}$ | Density-based Method (DBSCAN) | Dense sub-graphs |
| | FraudNE [136] | NR | $\mathcal{L}_{res}^{source} + \mathcal{L}_{res}^{sink} + \alpha \mathcal{L}_{share} + \eta \mathcal{L}_{reg}$ | Density-based Method (DBSCAN) | Dense sub-graphs |
| Anomalous Graph Detection Techniques | | | | | |
| Graph Database - Attributed | UPFD [141] | NR | $-(y \log(p) + (1-y) \log(1-p))$ | Anomaly Prediction | Predicted Label |
| | OCGIN [142] | GNN | $\min_W \frac{1}{N} \sum_{i=1}^N \ GIN(G_i, W) - c\ ^2 + \frac{\lambda}{2} \sum_{l=1}^L \ W^l\ _F^2$ | Location in Embedding Space | Distance to Hypersphere Center |
| Dynamic Graph - Plain | DeepSphere [147] | DNN | $\mathcal{L} = \mathcal{L}_h + \lambda \mathcal{L}_{res}$ | Location in Embedding Space | Anomalous Label |
| | GLAD-PAW [148] | GNN | cross-entropy($\mathbf{y}, \hat{\mathbf{y}}$) | Anomaly Prediction | Predicted Label |

* DNN: Deep NN Based Techniques, GCN: GCN Based Techniques, NR: Network Representation Based Techniques.

* GNN: Graph Neural Network Based Techniques.

idea is to seize the normal activity patterns and detect anomalies that behave significantly differently. The proposed method, OCAN, starts by extracting the benign users' content features using their historical social behaviors (e.g., historical posts, posts' URL), for which this method is classified into the dynamic category. A long short-term memory (LSTM) based autoencoder [125] is employed to achieve this and as assumed, benign users and malicious users are in separate regions in the feature space. Next, a novel one-class adversarial net comprising a generator and a discriminator is trained. Specifically, the generator produces complementary data points that locate in the relatively low density areas of benign users. The discriminator, accordingly, aims to distinguish the generated samples from the benign users. After training, benign users' regions are learned by the discriminator and anomalies can hence be identified with regard to their locations.

Both NetWalk [84] and OCAN [121] approach the anomalous node detection problem promisingly, however, they respectively only consider the structure or attributes. By the success of static graph anomaly detection techniques that analyze both aspects, when the structure and attribute information in dynamic graphs are jointly considered, an enhanced

detection performance can be foreseen. We therefore highlight this unexplored area for future works in Section 11.

5 ANOMALOUS EDGE DETECTION (ANOS ED)

In contrast to anomalous node detection, which targets individual nodes, ANOS ED aims to identify abnormal links. These links often inform the unexpected or unusual relationships between real objects [126], such as the abnormal interactions between fraudsters and benign users shown in Fig. 1, or suspicious interactions between attacker nodes and benign user machines in computer networks. Following the previous taxonomy, in this section, we review the state-of-the-art ANOS ED methods for static graphs, and Section 6 summarizes the techniques for dynamic graphs. A summary is provided in Table 3. This section includes methods based on deep NNs, GCNs and network representations. The non-deep learning techniques are reviewed in Appendix E, available in the online supplemental material.

5.1 Deep NN Based Techniques

Similar to deep NN based ANOS ND techniques, autoencoder and fully connected network (FCN) have also been

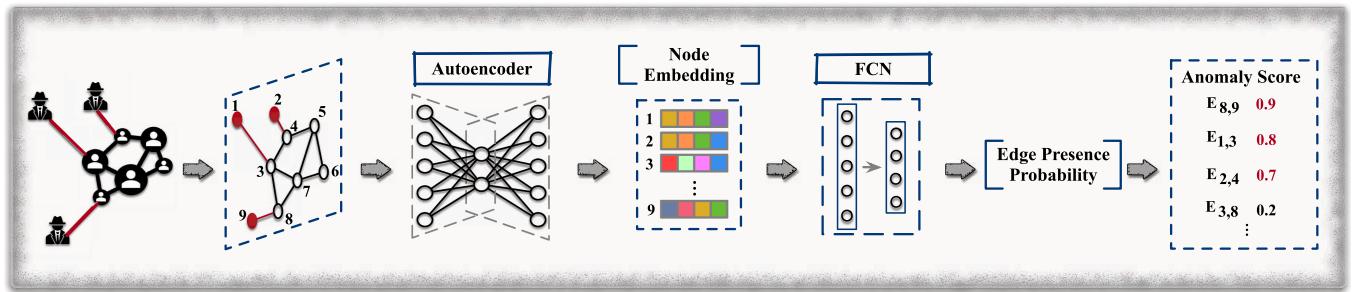


Fig. 6. ANOS ED on static graphs – Deep NN based approaches. For example, the detection technique employs the autoencoder and fully connected network to learn the presence probability of each edge. Anomaly scores are assigned with regard to the probabilities, and the top- k edges are anomalous (e.g., $E_{8,9}$, $E_{1,3}$, and $E_{2,4}$ at top-3).

used for anomalous edge detection. As an example, Ouyang *et al.* [127] approached the problem by modeling the distribution of edges through deep models to identify the existing edges that are least likely to appear as anomalies (as shown in Fig. 6). The probability of each edge (u, v) is decided by $P(v|u, N(u))$ and $P(u|v, N(v))$ which measure the edge probability using node u with its neighbors $N(u)$ and node v with its neighbors $N(v)$, respectively. To calculate $P(v|u, N(u))$, the proposed method, UGED, first encodes each node into a lower-dimensional vector through a FCN layer and generates node u 's representation by a mean aggregation of itself and its neighbors' vectors. Next, the node representations are fed into another FCN to estimate $P(v|u, N(u))$. The prediction is expressed as $\hat{P}(v|u, N(u)) = \text{Softmax}(W \cdot H(u))|_v$, where W represents the trainable parameters, and $H(u)$ is u 's representation. UGED's training scheme aims to maximize the prediction of existing edges via a cross-entropy-based loss function, $\text{CE}(\hat{P}(v|u, N(u)), v)$. After training, an anomaly score is assigned to each edge using the average of $1 - P(v|u, N(u))$ and $1 - P(u|v, N(v))$. As such, existing edges that have a lower probability will get higher scores and the top- k edges are reported as anomalous.

5.2 GCN Based Techniques

Following the line of modeling edge distributions, some studies leverage GCNs to better capture the graph structure information. Duan *et al.* [128] demonstrated that the existence of anomalous edges in the training data prevents traditional GCN based models from capturing real edge distributions, which leads to sub-optimal detection performance. This inherently raises a problem: to achieve better detection performance, the node embedding process should alleviate the negative impact of anomalous edges, but these edges are detected using the learned embeddings. To tackle this, the proposed method, AANE, jointly considers these two issues by iteratively updating the embeddings and detection results during training.

In each training iteration, AANE generates node embeddings Z through GCN layers and learns an indicator matrix I to spot potential anomalous edges. Given an input graph G with adjacency matrix A , each term I_{uv} in I is 1 if $\hat{A}_{uv} < \text{mean}_{v' \in N_u} \hat{A}_{uv'} - \mu \cdot \text{std}_{v' \in N_u} \hat{A}_{uv'}$, and 0 otherwise. Here, \hat{A}_{uv} is the predicted link probability between nodes u and v , which is calculated as the hyperbolic tangent of u and v 's embeddings, and μ is a predefined threshold. By this, an edge uv is identified as anomalous when its predicted

probability is less than the average of all links associated with the node u by a predefined threshold.

The total loss function of AANE contains two parts: an anomaly-aware loss (\mathcal{L}_{aal}) and an adjusted fitting loss (\mathcal{L}_{afl}). \mathcal{L}_{aal} is proposed to penalize the link prediction results and the indicator matrix I such that anomalous edges will have lower prediction probabilities when they are marked as 1 in I . This is formulated as:

$$\mathcal{L}_{aal} = \sqrt{\sum_{u \in V} \sum_{v \in N(u)} ((1 - \hat{A}_{uv}^2)(1 - I_{uv}) + \hat{A}_{uv}^2 I_{uv})}, \quad (11)$$

where V is the node set, $N(u)$ is the set of u 's neighbors. \mathcal{L}_{afl} quantifies the reconstruction loss with regard to the removal of potential anomalous edges, denoted as:

$$\mathcal{L}_{afl} = \|B - \hat{A}\|_2^2, \quad (12)$$

where B is an adjusted adjacency matrix that removes all predicted anomalies from the input adjacency matrix A . By minimizing these two losses, AANE identifies the top- k edges with lowest probabilities as anomalies.

5.3 Network Representation Based Techniques

Instead of using node embeddings for ANOS ED, edge representations learned directly from the graph are also feasible for distinguishing anomalies. If the edge representations well-preserve the graph structure and interaction content (e.g., messages in online social networks, co-authored papers in citation networks) between pairs of nodes, an enhanced detection performance can then be expected. To date, several studies, such as Xu *et al.* [129], have shown promising results in generating edge representations. Although they are not specifically designed for graph anomaly detection, they pinpoint a potential approach to ANOS ED. This is highlighted as a potential future direction in Section 11.1.

6 ANOS ED ON DYNAMIC GRAPHS

Dynamic graphs are powerful in reflecting the appearance/disappearance of edges over time [130]. Anomalous edges can be distinguished by modeling the changes in graph structure and capturing the edge distributions at each time step. Recent approaches to ANOS ED on dynamic graphs are reviewed in this section.

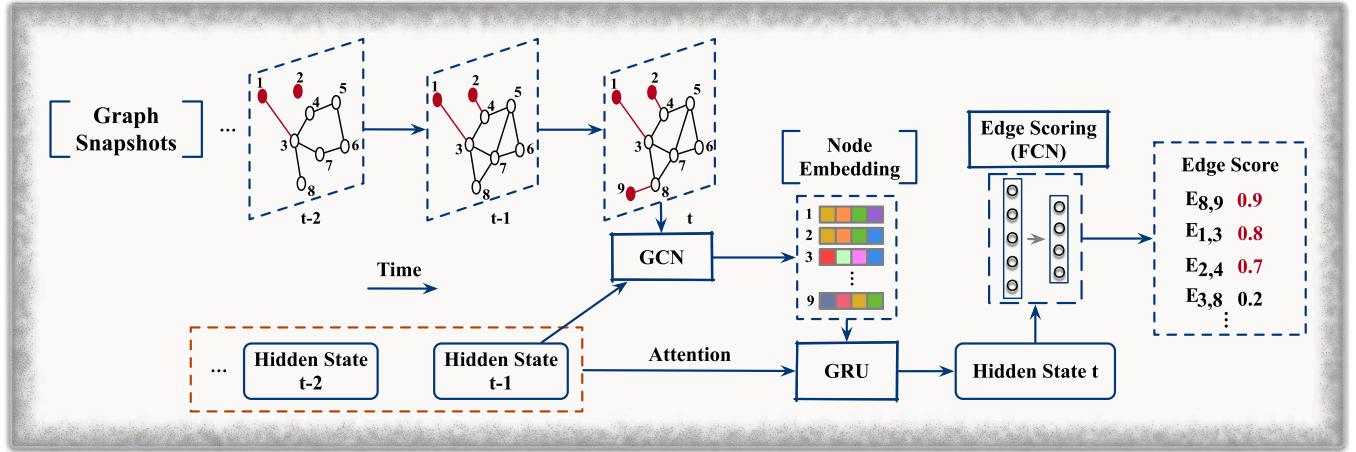


Fig. 7. ANOS ED on dynamic graphs – GCN based approaches. GCN is employed to learn node embeddings from the temporal graph at each timestamp. The attention-based GRU generates the current hidden state using the node embeddings and previous hidden states. The edge scoring function, such as a FCN, is learned to assign anomaly scores, and the top- k edges are depicted as anomalies.

6.1 Network Representation Based Techniques

The intuition of network representation based techniques is to encode the dynamic graph structure information into edge representations and apply the aforementioned traditional anomaly detection techniques to spot irregular edges. This is quite straightforward, but there remain vital challenges in generating/updating informative edge representations when the graph structure evolves. To mitigate this challenge, the ANOS ND model NetWalk [84] is also capable of detecting anomalous edges in dynamic graphs. Following the line of distance-based anomaly detection, NetWalk encodes edges into a shared latent space using node embeddings, and anomalies are identified based on their distances to the nearest edge-cluster centers in the latent space. Practically, Netwalk generates edge representations as the Hadamard product of the source and destination nodes' representations, denoted as: $\mathbf{z}_{u,v} = \mathbf{z}_u \odot \mathbf{z}_v$. When new edges arrive or existing edges disappear, the node and edge representations are updated from random walks in the temporary graphs at each time stamp, after which the edge-cluster centers and edge anomaly scores are recalculated. Finally, the top- k farthest edges to the edge-clusters are reported as anomalies.

6.2 GCN Based Techniques

Although NetWalk is capable of detecting anomalies in dynamic graphs, it simply updates edge representations without considering the evolving patterns of long/short-term nodes and the graph's structure. For more effective ANOS ED, Zheng *et al.* [131] intuitively combined temporal, structural and attribute information to measure the anomalousness of edges in dynamic graphs. They propose a semi-supervised model, AddGraph, which comprises a GCN and Gated Recurrent Units (GRU) with attention [132] to capture more representative structural information from the temporal graph in each time stamp and dependencies between them, respectively.

At each time stamp t , GCN takes the output hidden state (H^{t-1}) at time $t-1$ to generate node embeddings, after which the GRU learns the current hidden state H^t from the node embeddings and attentions on previous hidden states

(as shown in Fig. 7). After getting the hidden state H^t of all nodes, AddGraph assigns an anomaly score to each edge in the temporal graph based on the nodes associated with it. The proposed anomaly scoring function is formulated as:

$$f(u, v, w) = w \cdot \sigma(\beta \cdot (\|\mathbf{a} \odot \mathbf{h}_u + \mathbf{b} \odot \mathbf{h}_v\|_2^2 - \mu)), \quad (13)$$

where u and v are the corresponding nodes, w is the weight of the edge, \mathbf{a} and \mathbf{b} are trainable parameters, β and μ are hyper-parameters, and $\sigma(\cdot)$ is the non-linear activation function. To learn \mathbf{a} and \mathbf{b} , Zheng *et al.* further assumed that all existing edges in the dynamic graph are normal in the training stage, and sampled non-existing edges as anomalies. Specifically, they form the loss function as:

$$\begin{aligned} \mathcal{L}_{AddGraph} = \min & \sum_{(u, v, w) \in \varepsilon^t} \sum_{(u', v', w) \notin \varepsilon^t} \\ & \max\{0, \gamma + f(u, v, w) - f(u', v', w)\} + \lambda \mathcal{L}_{reg}, \end{aligned} \quad (14)$$

where ε^t is the edge set, (u', v') are sampled non-existing edges at time stamp t , λ is a hyper-parameter, and \mathcal{L}_{reg} regularizes all trainable parameters in the model. After training, the scoring function identifies anomalous edges in the test data by assigning higher anomaly scores to them based on Eq. (13).

7 ANOMALOUS SUB-GRAPH DETECTION (ANOS SGD)

In real life, anomalies might also collude and behave collectively with others to garner benefits. For instance, fraudulent user groups in an online review network, as shown in Fig. 1, may post misleading reviews to promote or besmirch certain merchandise. When these data are represented as graphs, anomalies and their interactions usually form suspicious sub-graphs, and ANOS SGD is proposed to distinguish them from the benign.

Unlike individual and independent graph anomalies, i.e., single nodes or edges, each node and edge in a suspicious sub-graph might be normal. However, when considered as a collection, they turn out to be anomalous. Moreover, these

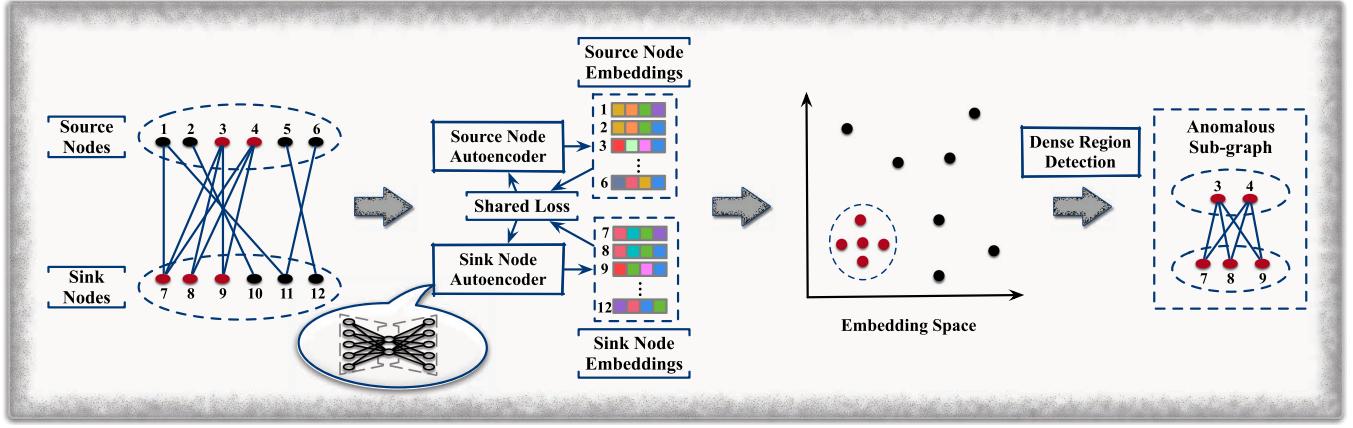


Fig. 8. ANOS SGD. Real-world networks are usually represented as bipartite graphs for reflecting the interactions between two different types of objects. To detect ANOS SGD, source nodes and sink nodes are embedded using two autoencoders (linked by a shared loss function), respectively. Anomalous sub-graphs are identified by applying dense region detection algorithms in the embedding space.

sub-graphs also vary in size and inner structure, making anomalous sub-graph detection more challenging than ANOS ND/ED [133]. Although extensive effort has been placed on circumventing this problem, deep-learning techniques have only begun to address this problem in the last five years. For reference, traditional non-deep learning based techniques are briefly introduced in Appendix F, available in the online supplemental material, and a summary of techniques reviewed for ANOS SGD is provided in Table 3.

Due to the flexibility of heterogeneous graphs in representing the complex relationships between different kinds of real objects, several recent works have taken advantage of deep network representation techniques to detect real-world anomalies through ANOS SGD. For instance, Wang *et al.* [134] represented online shopping networks as bipartite graphs (a specific type of heterogeneous graph that has two types of nodes and one type of edge), in which users are source nodes and items are sink nodes. Fraudulent groups are then detected based on suspicious dense blocks that form in these graphs.

Wang *et al.* [134] aimed to learn anomaly-aware representations of users such that suspicious users in the same group will be located closely in the vector space, while benign users will be far away (as shown in the embedding space in Fig. 8). According to the observation that user nodes belonging to one fraudulent group are more likely to connect with the same item nodes, the developed model, DeepFD, measures similarities in the behavior of two users, sim_{ij} , as the percentage of items shared among all the items they have reviewed. User representations are then generated through a traditional autoencoder, which is trained using three losses and follows the encoding-decoding process. The first loss is the reconstruction loss \mathcal{L}_{res} that ensures the bipartite graph structure can be reconstructed properly using the learned user representations and item representations. The second term \mathcal{L}_{sim} preserves the user similarity information in the learned user representations. That is, if two users have similar behaviors, their representations should also be similar. This loss is formulated as:

$$\mathcal{L}_{sim} = \sum_{i,j=1}^m sim_{ij} \cdot \|\widehat{sim}_{ij} - sim_{ij}\|_2^2, \quad (15)$$

where m is the number of user nodes, \widehat{sim}_{ij} measures the similarity of user i and j 's representations using an RBF kernel or other alternative. The third loss \mathcal{L}_{reg} regularizes all trainable parameters. Finally, the suspicious dense blocks, which are expected to form dense regions in the vector space, are detected using DBSCAN [135].

Another work, FraudNE [136], also models online review networks as bipartite graphs and further detects both malicious users and associated manipulated items following the dense block detection principle. Unlike DeepFD, FraudNE aspires to encode both types of nodes into a shared latent space where suspicious users and items belonging to the same dense block are very close to each other while others distribute uniformly. FraudNE adopts two traditional autoencoders, namely, a source node autoencoder and a sink node autoencoder, to learn user representations and item representations, respectively. Both autoencoders are trained to jointly minimize their corresponding reconstruction losses and a shared loss function, and the total loss can be formulated as:

$$\mathcal{L}_{FraudNE} = \mathcal{L}_{res}^{source} + \mathcal{L}_{res}^{sink} + \alpha \mathcal{L}_{share} + \eta \mathcal{L}_{reg}, \quad (16)$$

where α and η are hyperparameters, and \mathcal{L}_{reg} regularizes all trainable parameters. Specifically, the reconstruction losses (i.e., $\mathcal{L}_{res}^{source}$ and \mathcal{L}_{res}^{sink}) measure the gap between the input user/item features (extracted from the graph structure) and their decoded features. The shared loss function is proposed to restrict the representation learning process such that each linked pair of users and items get similar representations. As the DBSCAN [135] algorithm is convenient to apply for dense region detection, FraudNE also uses it to distinguish the dense sub-graphs formed by suspicious users and items.

To date, only a few works have put their efforts into using deep learning techniques for ANOS SGD. However, with intensifying research interest in sub-graph representation learning, we encourage more studies on ANOS SGD and highlight this as a potential future in Section 11.1.

8 ANOMALOUS GRAPH DETECTION (ANOS GD)

Beyond anomalous node, edge, and sub-graph, graph anomalies might also appear as abnormal graphs in a set/

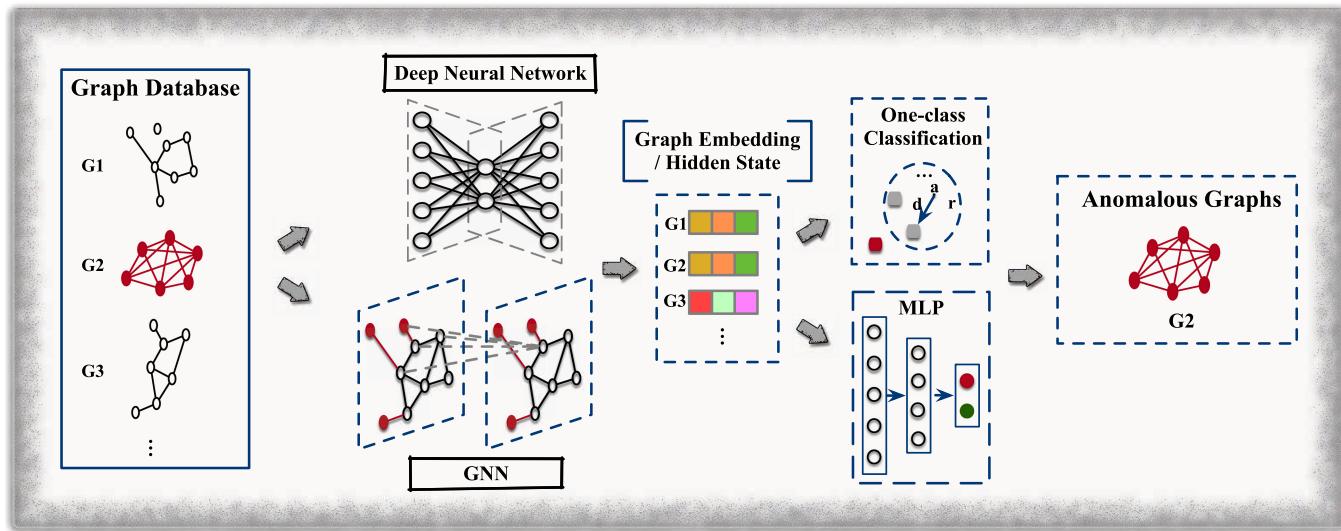


Fig. 9. ANOS GD on Graph Database. Generally, the graph-level anomaly detection techniques take a graph database as input. By generating embeddings/hidden state for each single graph in the database through D(G)NNs, anomalous graphs can be depicted by One-class classifiers or MLP.

database of graphs. Typically, a graph database is defined as:

Definition 4. (Graph Database). A graph database $\mathcal{G} = \{G_i = (V_i, E_i, X_v(i), X_e(i))\}_{i=1}^N$ contains N individual graphs. Here, each graph G_i is comprised of a node set V_i and an edge set E_i . $X_v(i)$ and $X_e(i)$ are the node attribute matrix and edge attribute matrix of G_i if it is an attributed graph.

This graph-level ANOS GD aims to detect individual graphs that deviate significantly from the others. A concrete example of ANOS GD is unusual molecule detection. When chemical compounds are represented as molecular/chemical graphs where the atoms and bonds are represented as nodes and edges [137], [138], unusual molecules can be identified because their corresponding graphs have structures and/or features that deviate from the others. Brain disorders detection is another example. A brain disorder can be diagnosed by analyzing the dynamics of brain graphs at different stages of aging in sequence and finding an inconsistent snapshot at a specific time stamp.

The prior reviewed techniques (i.e., ANOS ND/ED/SGD) are not compatible with ANOS GD because they are dedicated to detecting anomalies in a single graph, whereas ANOS GD is directed at detecting graph-level anomalies. This problem is commonly approached by: 1) measuring the pairwise proximities of graphs using graph kernels [139]; 2) detecting the appearance of anomalous graph signals created by abnormal groups of nodes [140]; or 3) encoding graphs using frequent motifs [63]. However, none of these methods are deep learning-based. As the time of writing, very few studies in ANOS GD with deep learning have been undertaken. As such, this is highlighted as a potential future direction in Section 11.1.

8.1 GNN Based Techniques

Motivated by the success of GNNs in various graph classification tasks, the most recent works in ANOS GD employ GNNs to classify single graphs as normal/abnormal in the given

graph database. Specifically, Dou *et al.* [141] transformed fake news detection into an ANOS GD problem by modeling news as tree-structured propagation graphs where the root nodes depict pieces of news, and child nodes depict users who interact with the root news. Their end-to-end framework, UPFD, extracts two embeddings for the news piece and users, respectively, via a text embedding model (e.g., word2vec, BERT) and a user engagement embedding process. For each news graph, its latent representation is a flattened concatenation of these two embeddings, which is input to train a neural classifier with the label of the news. Corresponding propagation graphs that are labeled as fake by the trained model are regarded as anomalous.

Another representative work by Zhao and Akoglu [142] employed a GIN model and one-class classification (i.e., DeepSVDD [143]) loss to train a graph-level anomaly detection framework in an end-to-end manner. For each individual graph in the graph database, its graph-level embedding is generated by applying mean-pooling over its nodes' node-level embeddings. A graph is eventually depicted as anomalous if it lies outside the learned hypersphere, as shown in Fig. 9.

8.2 Network Representation Based Techniques

It is also possible to apply general graph-level network representation techniques to ANOS GD. With these methods, the detection problem is transformed into a conventional outlier detection problem in the embedding space. In contrast to D(G)NN based techniques that can detect graph anomalies in an end-to-end manner, adopting these representation techniques for anomaly detection is two-staged. First, graphs in the database are encoded into a shared latent space using graph-level representation techniques, such as Graph2Vec [144], FGSD [145]. Then, the anomalousness of each single graph is measured by an off-the-shelf outlier detector. Essentially, this kind of approach involves pairing existing methods in both stages, yet, the stages are disconnected from each other and, hence, the detection performance can be subpar since the embedding similarities

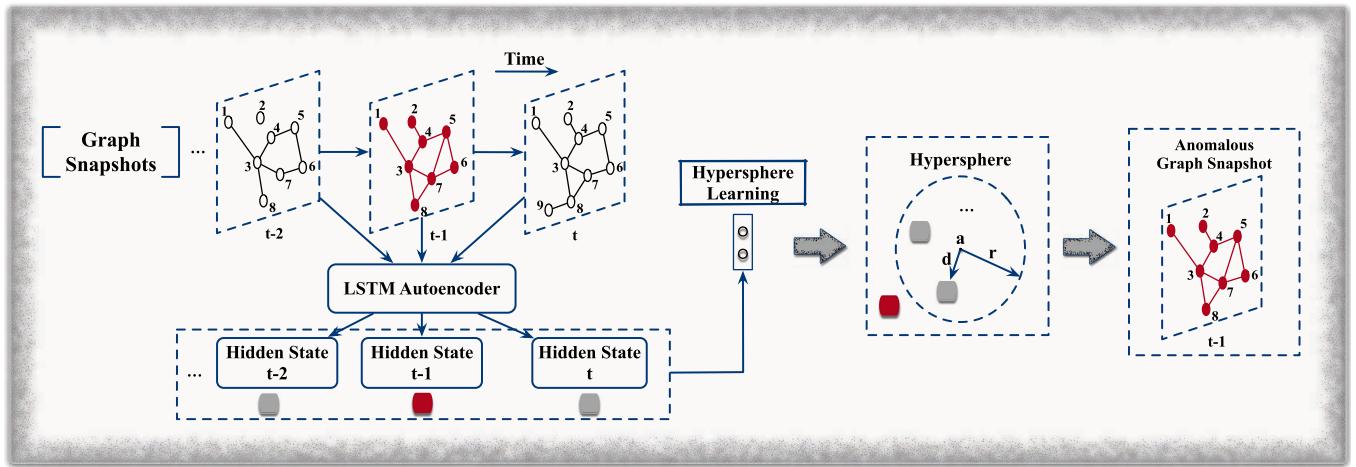


Fig. 10. ANOS GD on dynamic graphs. For each graph snapshot in the dynamic graph, the LSTM autoencoder generates its hidden state using its adjacency matrix and previous hidden states. Through hypersphere learning, a hypersphere with centriod a and radius r is learned such that anomalous snapshots lay outside.

are not necessarily designed for the sake of anomaly detection.

9 ANOS GD ON DYNAMIC GRAPHS

For dynamic graph environments, graph-level anomaly detection endeavors to identify abnormal graph snapshots/temporal graphs. Similar to ANOS ND and ED on dynamic graphs, given a sequence of graphs, anomalous graphs can be distinguished regarding their unusual evolving patterns, abnormal graph-level features, or other characteristics.

In order to derive each graph snapshot/temporal graph's characteristics, the commonly used GNN, LSTM and autoencoder are feasible to apply. For instance, Teng *et al.* [147] applied a LSTM-autoencoder to detect abnormal graph snapshots, as shown in Fig. 10. In their proposed model, DeepSphere, a dynamic graph is described as a collection of three-order tensors, $\{\mathcal{X}_k, k = 1, 2, \dots\}$ where each $\mathcal{X} \in \mathbb{R}^{N \times N \times T}$, and the slices along the time dimension are the adjacency matrices of graph snapshots. To identify abnormal tensors, DeepSphere first embeds each graph snapshot into a latent space using an LSTM autoencoder, and then leverages a one-class classification objective [143] that learns a hypersphere such that normal snapshots are covered, and anomalous snapshots lay outside. The LSTM autoencoder takes the adjacency matrices as input sequentially and attempts to reconstruct these input matrices through training. And the hypersphere is learned through a single neural network layer and its objective function is formulated as:

$$\mathcal{L}_h = r^2 + \gamma \sum_{k=1}^m \epsilon_k + \frac{1}{m} \sum_{k=1}^m \|\mathbf{z}_k - \mathbf{a}\|^2, \quad (17)$$

where \mathbf{z}_k is the latent representation generated by the LSTM autoencoder, \mathbf{a} is the centroid of the hypersphere, r is the radius, ϵ_k is the outlier penalty ($\epsilon_k = \|\mathbf{z}_k - \mathbf{a}\|^2 - r^2$), m is the number of training graph snapshots, and γ is a hyper-parameter. The overall objective function of DeepSphere is represented as:

Authorized licensed use limited to: Anhui University. Downloaded on June 03, 2025 at 02:34:12 UTC from IEEE Xplore. Restrictions apply.

$$\mathcal{L} = \mathcal{L}_h + \lambda \mathcal{L}_{res}, \quad (18)$$

where \mathcal{L}_{res} is the reconstruction loss of the LSTM autoencoder. When the training is finished, DeepSphere spots a given unseen data \mathcal{X} as anomalous if its embedding lies outside the learned hypersphere with a radius of r .

In addition to all ANOS ND, ED, SGD, and GD techniques reviewed above, it is worth mentioning that perturbed graphs, which adversarial models generate to attack graph classification algorithms or GNNs [149], [150], [151], can also be regarded as (intensional) anomalies. In a perturbed graph, the nodes and edges are modified deliberately to deviate from the others. We have not reviewed these in this survey because their main purpose is to attack a GNN model. The key idea behind these methods is the attacking/perturbation strategy, and studies in this sphere seldom focus on a detection or reasoning module to identify the perturbed graph or its sub-structures, i.e., anomalous nodes, edges, sub-graphs, or graphs.

10 PUBLISHED ALGORITHMS AND DATASETS

Acquiring open-sourced implementations and real-world datasets with real-world anomalies are far from trivial in academic research on graph anomaly detection. Here, we first list the published algorithms with publicly available implementations, then we provide a collection of public benchmark datasets and summarize the commonly used evaluation metrics. Lastly, due to the shortage of labeled anomalies in real-world datasets, we review three synthetic dataset generation strategies used in the existing works. All the resources are available at: <https://xiaoxiaoma-mq.github.io/Awesome-Deep-Graph-Anomaly-Detection/>.

10.1 Published Algorithms

The published implementations of algorithms and models contribute to baseline experiments. Table 4 provides a summary of published implementations outlining the language and platforms, graphs that they can admit, and URLs to code repositories.

10.2 Published Datasets

Table 5 summarizes the most commonly used datasets, categorizing them into different groups with regard to their

TABLE 4
Published Algorithms and Models

| Model | Language | Platform | Graph | Code Repository |
|-------------------|----------|------------|-------------------------|---|
| AnomalyDAE [110] | Python | Tensorflow | Static Attributed Graph | https://github.com/haoyfan/AnomalyDAE |
| MADAN [153] | Python | - | Static Attributed Graph | https://github.com/leoguti85/MADAN |
| PAICAN [72] | Python | Tensorflow | Static Attributed Graph | http://www.kdd.in.tum.de/PAICAN/ |
| ONE [83] | Python | - | Static Attributed Graph | https://github.com/sambaranban/ONE |
| DONE&AdONE [82] | Python | Tensorflow | Static Attributed Graph | https://bit.ly/35A2xHs |
| SLICENDICE [154] | Python | - | Static Attributed Graph | http://github.com/hamedn/SliceNDice/ |
| FRAUDRE [118] | Python | Pytorch | Static Attributed Graph | https://github.com/FraudDetection/FRAUDRE |
| SemiGNN [111] | Python | Tensorflow | Static Attributed Graph | https://github.com/safe-graph/DGFraud |
| CARE-GNN [27] | Python | Pytorch | Static Attributed Graph | https://github.com/YingtongDou/CARE-GNN |
| GraphConsis [155] | Python | Tensorflow | Static Attributed Graph | https://github.com/safe-graph/DGFraud |
| GLOD [142] | Python | Pytorch | Static Attributed Graph | https://github.com/LingxiaoShawn/GLOD-Issues |
| OCAN [121] | Python | Tensorflow | Static Graph | https://github.com/PanpanZheng/OCAN |
| DeFrauder [166] | Python | - | Static Graph | https://github.com/LCS2-IIITD/DeFrauder |
| GCAN [156] | Python | Keras | Heterogeneous Graph | https://github.com/l852888/GCAN |
| HGATRD [157] | Python | Pytorch | Heterogeneous Graph | https://github.com/201518018629031/HGATRD |
| GLAN [158] | Python | Pytorch | Heterogeneous Graph | https://github.com/chunyuanY/RumorDetection |
| GEM [164] | Python | - | Heterogeneous Graph | https://github.com/safe-graph/DGFraud/tree/master/algorithms/GEM |
| eFraudCom [146] | Python | Pytorch | Heterogeneous Graph | https://github.com/GeZhangMQ/eFraudCom |
| DeepFD [134] | Python | Pytorch | Bipartite Graph | https://github.com/JiaWu-Repository/DeepFD-pyTorch |
| ANOMRANK [159] | C++ | - | Dynamic Graph | https://github.com/minjiyoon/anomrank |
| MIDAS [165] | C++ | - | Dynamic Graph | https://github.com/Stream-AD/MIDAS |
| Sedanspot [152] | C++ | - | Dynamic Graph | https://www.github.com/dhiyyaeswaran/sedanspot |
| F-FADE [126] | Python | Pytorch | Dynamic Graph | http://snap.stanford.edu/f-fade/ |
| DeepSphere [147] | Python | Tensorflow | Dynamic Graph | https://github.com/picsolab/DeepSphere |
| Changedar [140] | Matlab | - | Dynamic Graph | https://bhooi.github.io/changedar/ |
| UPFD [141] | Python | Pytorch | Graph Database | https://github.com/safe-graph/GNN-FakeNews |
| OCGIN [142] | Python | Pytorch | Graph Database | https://github.com/LingxiaoShawn/GLOD-Issues |
| DAGMM [160] | Python | Pytorch | Non Graph | https://github.com/danieltan07/dagmm |
| DevNet [161] | Python | Tensorflow | Non Graph | https://github.com/GuansongPang/deviation-network |
| RDA [162] | Python | Tensorflow | Non Graph | https://github.com/zc8340311/RobustAutoencoder |
| GAD [163] | Python | Tensorflow | Non Graph | https://github.com/raghavchhalapathy/gad |
| Deep SAD [168] | Python | Pytorch | Non Graph | https://github.com/lukasruff/Deep-SAD-PyTorch |
| DATE [169] | Python | Pytorch | Non Graph | https://github.com/Roysai27/Dual-Attentive-Tree-aware-Embedding |
| STS-NN [167] | Python | Pytorch | Non Graph | https://github.com/JiaWu-Repository/STS-NN |

* -: No Dedicated Platforms.

application fields. Notably, there is a lack of ground truth with labeled anomalies only provided in the Enron, Twitter Sybil, Disney, Amazon, Elliptic and Yelp datasets. Details of the DBLP, UCI message, Digg, Wikipedia, and New York city taxi datasets are not given because these public datasets only contain the raw data, and in most existing works, they are further processed to build different graphs (e.g., homogeneous graphs, bipartite graphs). The well-known citation

networks are often used to generate synthetic datasets by injecting anomalies into them - the number of anomalies varies from study to study.

In addition to these anomaly detection datasets, Table 5 also lists eight graph classification datasets. As mentioned in 10.3, through downsampling, these datasets can be used as benchmarks for evaluating the anomaly detection performance.

TABLE 5
Published Datasets

| Category | Dataset | #G | #N | #E | #FT | #AN | REF | URL |
|-------------------------|-----------------------|------|------|------|------|------|---|---|
| Citation Networks | ACM | 1 | 16K | 71K | 8.3K | - | [92], [100], [110], [112] | http://www.arnetminer.org/open-academic-graph |
| | Cora | 1 | 2.7K | 5.2K | 1.4K | - | [72], [82], [83], [97], [114] | http://linqs.cs.umd.edu/projects/projects/lbc |
| | Citeseer | 1 | 3.3K | 4.7K | 3.7K | - | [12], [82], [83], [114] | http://linqs.cs.umd.edu/projects/projects/lbc |
| | Pubmed | 1 | 19K | 44K | 500 | - | [82], [83], [97], [114] | http://linqs.cs.umd.edu/projects/projects/lbc |
| | DBLP | 1 | - | - | - | - | [12], [72], [77], [84], [152] | http://www.informatik.uni-trier.de/~ley/db/ |
| Social Networks | Enron | - | 80K | - | - | - | [39], [41], [104], [107], [113], [152], [153], [159], [170] | http://odds.cs.stonybrook.edu/#table2 |
| | UCI Message | 1 | 5K | - | - | - | [84], [85], [131] | http://archive.ics.uci.edu/ml |
| | Google+ | 4 | 75M | 11G | - | - | - | https://wangbinghui.net/dataset.html |
| | Twitter Sybil | 3 | 41M | - | - | 100K | - | https://wangbinghui.net/dataset.html |
| | Twitter World-Cup2014 | - | 54K | - | - | - | [170] | SelectiveAnomalyEnsemble/">http://shebuti.com>SelectiveAnomalyEnsemble/ |
| | Twitter Security2014 | - | 130K | - | - | - | [170] | SelectiveAnomalyEnsemble/">http://shebuti.com>SelectiveAnomalyEnsemble/ |
| | Reality Mining | - | 9.1K | - | - | - | [170] | SelectiveAnomalyEnsemble/">http://shebuti.com>SelectiveAnomalyEnsemble/ |
| | NYTNews | - | 320K | - | - | - | [170] | SelectiveAnomalyEnsemble/">http://shebuti.com>SelectiveAnomalyEnsemble/ |
| | PolitiFact | 314 | 41K | 40K | - | 157 | [141] | https://github.com/safe-graph/GNN-FakeNews |
| | Gossipcop | 5.4K | 314K | 308K | - | 2.7K | [141] | https://github.com/safe-graph/GNN-FakeNews |
| Co-purchasing Networks | Disney | 1 | 124 | 334 | 30 | 6 | [41], [106], [107], [113], [153] | https://www.ipd.kit.edu/mitarbeiter/muellere/consub/ |
| | Amazon-v1 | 1 | 314K | 882K | 28 | 6.2K | [8], [26], [72], [73], [76], [77], [107] | https://www.ipd.kit.edu/mitarbeiter/muellere/consub/ |
| | Amazon-v2 | 1 | 11K | - | 25 | 821 | - | https://github.com/dmlc/dgl/blob/master/python/dgl/data/fraud.py |
| | Elliptic | 1 | 203K | 234K | 166 | 4.5K | - | https://www.kaggle.com/ellipticco/elliptic-data-set |
| Transportation Networks | Yelp | 1 | 45K | - | 32 | 6.6K | - | https://github.com/dmlc/dgl/blob/master/python/dgl/data/fraud.py |
| | New York City Taxi | - | - | - | - | - | [39], [120], [147] | http://www.nyc.gov/html/tlc/html/about/triprecorddata.shtml |

* -: Not Given, #G: Number of Graphs, #N: Number of Nodes, #E: Number of Edges, #FT: Number of Features, #AN: Number of Anomalies, REF: References.

10.3 Synthetic Dataset Generation

Given the rarity of ground-truth anomalies, many researchers have employed synthetic datasets to investigate the effectiveness of their proposed methods [83], [171], [172]. Typically, these datasets can be categorized as follows:

- *Synthetic graphs with injected anomalies.* Pursuing this strategy, graphs are created to simulate real-world networks. All the nodes and edges are manually added with well-known benchmarks (e.g., Lanchinetti-Fornunato-Radicchi (LFR) [173], small-world [174], scale-free graphs [175]). Once built, ground-truth anomalies are planted into the network. For the feasibility of generating expected scale of networks, this strategy is mostly used by previous works to validate their underlying intuitions in anomaly detection.
- *Real-world datasets with injected anomalies.* These datasets are built based on the real-world networks. In particular, anomalies are created either by modifying the topological structure or the attributes of existing nodes/edges/sub-graphs, or by inserting non-existent graph objects.
- *Downsampled graph classification datasets.* The widely-used graph classification datasets (e.g., NCI1, IMDB, ENZYMES in [142]) can be easily converted into sets

suitable for anomaly detection through two steps. First, one particular class and its data records are chosen to represent normal objects. Then, the other data records are downsampled as anomalies at a specified downsampling rate. By this, the generated graph anomaly detection dataset is, in fact, a subset of the original dataset. The most significant strength of this strategy is that no single data record has been modified.

10.4 Evaluation Metrics

To date, the most widely used metrics for evaluating anomaly detection performance evaluation include accuracy, precision, recall rate, F1-score and AUC-AP (Average Precision). Their formulas/descriptions are given in Table 6. However, a more dedicated analysis with some new evaluation metrics is needed for further performance examination because anomaly detection has different requirements for different applications [176], [177], [178], e.g., false negative and false positive. For instance, network intrusion prevention systems are more sensitive to false negative errors, while false positive errors are considered relatively harmless. This is because any risky connections should be shut down to prevent information leaks. By contrast, other applications concentrate more on the false positives, e.g., in auditing domain, companies often set a budget

TABLE 6
Evaluation Metrics

| Evaluation Metric | Formula/Description |
|-------------------|---|
| Accuracy | $\frac{tp+tn}{tp+tn+fp+fn}$ |
| Precision | $\frac{tp}{tp+fp}$ |
| Recall | $\frac{tp}{tp+fn}$ |
| F1 Score | $2 * \frac{Recall * Precision}{Recall + Precision}$ |
| AUC-ROC | The Area Under the ROC Curve |
| AUC-AP | The Area Under Precision-Recall Curve |

tp: true positives; *tn*: true negatives; *fp*: false positives; *fn*: false negatives.

for an auditor to look at flagged anomalies, and they want high precision/small false positive rate such that the auditor's time can be best used. Hence, when evaluating detection performance, we suggest reviewing the specific requirements of the application domain for fair and suitable comparisons.

11 FUTURE DIRECTIONS

So far, we have reviewed the contemporary deep learning techniques that are devoted to graph anomaly detection. An apparent observation from our survey is there remain many compounded challenges imposed by the complexity of anomaly detection, graph data, and the immaturity of deep learning techniques for graph data mining. Another observation is that deep learning techniques in graph anomaly detection are still confined to a relatively small number of studies, and most of these focus on anomalous node detection. For a gauge, simply compare the length of Tables 2 and 3. Edge, sub-graph, and graph-level anomaly detection have clearly received much less attention. To bridge the gaps and push forward future work, we have identified 12 directions of future research for graph anomaly detection with deep learning.

11.1 Anomalous Edge, Sub-Graph, and Graph Detection

In real-world graphs, anomalies also appear as unusual relationships between objects, sub-structures formed by abnormal groups, or abnormal graphs, which are known as anomalous edges, sub-graphs, and graphs respectively. As indicated in our review, there is a huge gap between the existing anomalous edge/sub-graph/graph detection techniques and the emerging demands for more advanced solutions in various application domains (e.g., social networks, computer networks, financial networks). When detecting anomalous edges/sub-graphs/graphs, the proposed methods should be capable of leveraging the rich information contained in graphs to find clues and characteristics that can distinguish normal objects and anomalies in specific applications. Typically, this involves extracting edge/sub-graph/graph-level features, modeling the patterns of these features, and measuring the abnormalities accordingly. However, current deep learning based graph anomaly detection techniques put forward very little effort in this regard.

Opportunities: We believe more research efforts can be done on anomalous edge, sub-graph, and graph detection with regard to their significance in real-world applications. Possible solutions to this gap might to be first consider the application domain and explore domain knowledge to find complementary clues as a basis for these problems. Then,

motivated by recent advances in deep learning for edge, sub-graph, and graph-level representation learning [129], [179], extensive work can be done to learn an anomaly-aware embedding space such that it is feasible to extract abnormal patterns of anomalies. Although this direction seems quite straightforward, the true challenge lies in the specific application domains. Hence, domain knowledge, anomalous pattern recognition and anomaly-aware deep learning techniques should be enforced simultaneously.

11.2 Anomaly Detection in Dynamic Graphs

Dynamic graphs provide powerful machinery with which to capture the evolving relationships between real objects and their attributes. Their ever-changing structure and attribute information inherently make anomaly detection very challenging in these scenarios, leading to two primary concerns for the task. The first is to consider the spatial and temporal information contained in each graph snapshot at different time stamps, and the second is to explore the evolving patterns of nodes, edges, sub-graphs and graphs, as well as their interaction with the node/edge attributes over time. When these challenges have been tackled with mature solutions, detection techniques will achieve better results.

Opportunities: From our observations, most of deep learning based dynamic graph anomaly detection techniques are built on DeepWalk [79], GCN [91] or other deep models that are intuitively designed for static graphs. This means other information, like evolving patterns in attributes [180], [181], are not adequately used in the detection task. We can therefore identify the following directions for future studies to target.

- *Using dynamic graph mining tools.* As a popular research topic, deep learning for dynamic graph data mining [182], [183] has shown its effectiveness in supporting many dynamic graph analysis tasks, such as node clustering and graph classification [60]. More future works can be foreseen that adopt these techniques for anomaly detection.
- *Deriving solid evidence for anomaly detection.* The rich structural, attribute and temporal information in dynamic graphs are valuable resources for identifying anomalies. Apart from the indicators widely used in current works, such as burst of connections between node pairs or suddenly vanishing connections, we suggest exploring structural and attribute changes in depth. From such studies, we may derive additional information to enhance the detection performance, such as the appearance of abnormal attributes.
- *Handling complex dynamics.* Real-world networks always exhibit changes in both the network structure and node attributes, but only very few studies address this circumstance. Most of the state-of-arts only consider changes in one of these aspects. Although this 'double' scenario is extremely complex and detecting anomalies in this kind of dynamic graph is very challenging, it is worth studying because these graphs are highly reflective of real network data.

11.3 Anomaly Detection in Heterogeneous Graphs

Heterogeneous graphs are a specific type of graph that contain various types of nodes and edges. For instance, Twitter

can be intuitively modeled as a heterogeneous graph comprised of tweets, users, words, etc.

Opportunities: To use the complex relationships between different types of nodes in heterogeneous graphs for anomaly detection, representative works, such as HGATRD [157], GCAN [156] and GLAN [158], typically decompose a heterogeneous graph into individual graphs according to meta-paths, e.g., one with tweets and users, and another with tweets and words. They then use D(G)NNs to learn the embeddings for graph anomaly detection. Such a decomposition inherently overlooks the direct relations between different types of nodes/edges and downgrades the effectiveness of the embeddings. A possible solution is to reveal the complex relations between different types of nodes and edges, and encode them into a unique representation for boosted detection performance.

11.4 Anomaly Detection in Large-Scale Graphs

The scalability of methods to high-dimensional and large-scale data is an ongoing and significant challenge to anomaly detection techniques. In face of large-scale networks, such as Facebook and Twitter that contain billions of users and friendship links, the size of data in terms of both graph size and number of node attributes is extremely high. However, most of the existing works lack the ability to detect anomalies in such large-scale data because they are transductive models and need to take the whole graph as input for further analysis. Computation time and memory cost increase dramatically as the network scales up, and this stops existing techniques from being used on large-scale networks.

Opportunities: Accordingly, there is a need for scalable graph anomaly detection techniques. One possible approach would be an inductive learning scheme that first trains a detection model on part of the whole graph and then applies the model to detect anomalies in the unseen data. As some inductive learning models, such as GraphSAGE [184], have shown their effectiveness on link prediction and node classification in large-scale graphs, this approach is expected to provide a basis for graph anomaly detection in large-scale graphs and similar techniques can be investigated in the future.

11.5 Multi-View Graph Anomaly Detection

In real-world networks, objects might form different kinds of relationships with others (e.g., user's followership and friendship on Twiter). And their attribute information might be collected from different resources, such as user's profile, historical posts. This results in two types of multi-view graphs: 1) multi-graph that contains more than one type of edges between two nodes [185], [186]; and 2) multi-attributed-view graph that stores node attributes in different attributed views [93], [187], [188].

Opportunities: These multi-views basically allow us to analyze real objects' characteristics from different perspectives. Each view also provides complementary information to other views, and they might have different significance on anomaly detection. For instance, anomalies might be indistinguishable in one view but are obviously divergent from the majority in another view. There are a variety of work in data mining on multi-view learning [189], [190]. However, work that can accommodate multi-view graphs

along with multi-view attributes on nodes for anomaly detection purposes is nascent. Moreover, the rich information contained in multiple views and the inconsistency among them are overlooked in these works. To this end, we believe more research effort in this direction is needed. Digesting the relationships between views will be vital to their success, as two views might provide contrary/supplementary information for anomaly detection.

11.6 Camouflaged/Adversarial Anomaly Detection

The easy accessibility of online platforms has made them convenient targets for fraudsters, attackers and other malevolent agents to carry out malicious activities. Although various anomaly detection systems have been deployed to protect benign objects, anomalies can still conceal themselves to evade detection [191]. Known as camouflaged anomalies, these entities typically disguise themselves as regular objects. If the detection techniques are not robust against such cases, i.e., if they cannot quickly and effectively adapt to the evolving behavior of evasion-seeking attackers, the anomalies are simply left to cause their damage.

Opportunities: In the face of camouflage, the boundary between anomalies and regular objects is blurred, making anomalies much harder to be identified. We believe extensive effort should be placed on detecting these anomalies because, as yet, very few studies have looked at handling camouflaged anomalies in graphs [26], [27], [76]. To fulfill this gap, one major direction might be to jointly analyze the attributes, co-relations, such as the triadic, tetradic, or high-order relationships between objects in hypergraphs [192], [193], [194], [195], and other information comprised in graphs. By this, anomalies that only camouflage their local structures or attributes can be identified effectively. Enhancing existing techniques might be another direction. This involves incorporating additional detection mechanisms or function blocks particularly designed for distinguishing camouflaged anomalies with existing detection techniques. Consequently, these techniques will bridge most existing works and camouflaged anomaly detection.

11.7 Imbalanced Graph Anomaly Detection

Anomalies are rare, which means anomaly detection always occurs coexists with class imbalance in the training data. As deep learning models rely heavily on the training data, such imbalance pose great challenges to graph anomaly detection, and this remains a significant obstacle to deep learning techniques. Typically, the imbalanced class distributions will downgrade the detection techniques' ability to capture the differences between anomalies and non-anomalies. It might even cause over-fitting on the anomalous class because there are too few anomalies in the data. If the detection model overlooks this critical fact and is trained improperly, detection performance will be sub-optimal.

Opportunities: In fact, class imbalance has been widely explored in various research areas [117], [118]. Advances such as under-sampling the majority class or modifying the algorithms shed important light on solving training problems with imbalances. Yet, contemporary graph anomaly detection methods rarely incorporate these techniques. For more effective detection techniques, biased models that pay

more attention to anomalies, such as penalizing additional training losses on misclassified anomalies, would be a possible direction to circumvent the problem. Moreover, when adopting graph neural networks that aggregate neighboring information to the target node, like GCN or GraphSAGE, the over-smoothing between the features of connected nodes should be prevented such that the distinguishable features of the anomalies can be preserved to support anomaly detection.

11.8 Multi-Task Anomaly Detection

Graph anomaly detection has close relations with other graph mining tasks including community detection [58], node classification [196], and link prediction [197]. For a concrete example, when detecting community anomalies, community detection techniques are usually used to extract the community structures prior to anomaly detection. Meanwhile, the anomaly detection results can be used to optimize the community structure. Such mutually beneficial collaborations between anomaly detection and other tasks inherently suggest an opportunity for multi-task learning that can handle diverse tasks simultaneously and share information among tasks.

Opportunities: Multi-task learning provides effective machinery with which incorporate associated tasks [198], [199]. Its utmost advantage is that the training signal from another task could yield complementary information to distinguish anomalies from non-anomalies. The result would be enhanced detection performance. However, very few attempts focus on this at present. Beyond current works, such as [99] that jointly perform anomalous node detection and personalized recommendation, explorations into combining other learning tasks with graph anomaly detection are likely to emerge as a fruitful future direction.

11.9 Graph Anomaly Interpretability

The interpretability of anomaly detection techniques is vital to the subsequent anomaly handling process. When applying these techniques to real applications, such as financial and insurance systems, it is essential to provide explainable and lawful evidence to support the detection results. However, most of the existing works lack the ability to provide such evidence. To identify the anomalies, the most commonly used metrics are top- k rankings and simple anomaly scoring functions. These metrics are flexible enough to label objects as being either an anomaly or not an anomaly, but they cannot derive solid explanations. Moreover, as deep learning techniques have also been criticized for their low interpretability, future works on graph anomaly detection with deep learning should pay much more attention to this [200].

Opportunities: To bridge this gap, integrating specially designed interpretation algorithms or mechanisms [201], [202] into the detection framework would be a possible solution, noting that this would inherently induce a higher computational cost. Future works should therefore balance the cost of anomaly detection performance and interpretability. Visualization-based approaches, e.g., dashboards, charts, might also be feasible for showing the distinction between anomalies and non-anomalies in a human-friendly

manner. Further research in this direction will be successful if interpretable visualization results can be given [203].

11.10 Graph Anomaly Identification Strategies

Amongst existing unsupervised graph anomaly detection techniques, anomalies are mainly identified based on residual analysis [41], [107], reconstruction loss [110], distance-based statistics [84], density-based statistics [97], graph scan statistics [204], [205], [206], [207], and one-class classification [115]. The underlying intuition of these identification strategies is that anomalies have inconsistent data patterns with regular objects, and they will, therefore: 1) introduce more residual errors or be harder to reconstruct; or 2) be located in low-density areas or far away from the majority class in an anomaly-aware feature space. Effort toward designing novel loss functions for GNNs for anomaly detection is currently quite limited [67].

Opportunities: Although these strategies could capture the deviating data patterns of anomalies, they also have different limitations. Specifically, the residual analysis, one-class classification and reconstruction loss strategies are sensitive to noisy training data. Noisy nodes, edges or sub-graphs also exhibit large residuals, a large distance to the origin/hyper-sphere center and high reconstruction losses. Meanwhile, the distance-based and density-based strategies can only be applied when anomalies and non-anomalies are well separated in lower-dimensional space. Detection performance also downgrades dramatically if the gap between anomalies and non-anomalies is not that evident. It calls for extensive future efforts to break these limitations and explore new anomaly identification strategies.

11.11 Systematic Benchmarking

Systematic benchmarking is key to evaluating the performance of graph anomaly detection techniques. As indicated in our analysis in Section 10.4, recent studies are constantly raising attention to more comprehensive and effective benchmarking [142], [176], [177], [178]. Typically, the benchmarking framework consists of benchmark datasets, baseline methods, evaluation metrics, and further analysis tools. When evaluating a techniques' performance with other baselines, the evaluation dataset and metrics become very important because the performance of each model may vary depending on the setting. The shortage of public datasets and (public available) baseline methods also imposes great challenges for effective evaluation. Although one of the aims of our survey is to provide extensive materials for this purpose, like open-sourced implementations, datasets, and evaluation metrics, this work can only serve as a basis for future investigations into systematic benchmarking. We invite more efforts from the anomaly detection community toward this important case. Certainly, rigorous attention to designing better frameworks for benchmarking would help to disclose the advances and shortcomings of various detection techniques and essentially track a unprejudiced and accurate progress record in this field.

11.12 Unified Anomaly Detection Framework

A graph anomaly can be categorized as an anomalous node, edge, and sub-graph in a single graph or an anomalous

graph in a graph database. These anomalies usually coexist in real-world datasets. For instance, individual fraudsters, abnormal relationships, and fraud groups exist concurrently in online social network, as shown in Fig. 1. Moreover, there may be different ways to define anomalies of a certain type, such as community outliers versus anomalous communities or attribute-based versus structural anomalies. When deploying detection techniques in real applications, it is expected that all types of anomalies can be identified while consuming the least resources and time. A straightforward approach would be to integrate independent anomalous node, edge, and sub-graph detection techniques. Although this is convenient to apply to relatively small networks, its high computational cost will surely prevent the approach from scaling to large networks, such as Facebook and Twitter, because the same graph data has to be loaded and processed more than once by different techniques.

Opportunities: Unified frameworks that can detect diverse types of anomalies together [208], [209] may provide feasible solutions to bridge the gap. To build such frameworks, one possible direction is to capture all the information needed by different detection techniques simultaneously so that these techniques can be applied. The idea seems to be non-challenging, but in deep learning, methods of designing neural network layers and learning strategies that can fulfill this need will require extensive effort.

12 CONCLUSION

Due to the complex relationships between real-world objects and recent advances in deep learning, especially graph neural networks, graph anomaly detection with deep learning is currently at the forefront of anomaly detection. To the best of our knowledge, this is the first survey to present a comprehensive review dedicated to graph anomaly detection with modern deep learning techniques. Specifically, we have reviewed and categorized contemporary deep learning techniques according to the types of graph anomalies they can detect, including: (1) anomalous node detection; (2) anomalous edge detection; (3) anomalous sub-graph detection; and, finally, (4) anomalous graph detection. Clear summarizations and comparisons between different works are given to offer a complete and thorough picture of the current work and progress of graph anomaly detection as a field.

Moreover, to push forward future research in this area, we have provided a basis for systematic benchmarking by compiling a wide range of commonly used datasets, open-sourced implementations and synthetic dataset generation techniques. We further highlight 12 potential directions for future work based on the survey results. It is our firm belief that graph anomaly detection with deep learning is indeed more than a burst of temporary interest, and numerous applications from diverse domains are to surely benefit from it for the years to come.

REFERENCES

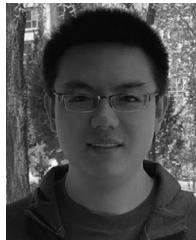
- [1] F. E. Grubbs, "Procedures for detecting outlying observations in samples," *Technometrics*, vol. 11, no. 1, pp. 1–21, 1969.
- [2] B. Hooi *et al.*, "Birdnest: Bayesian inference for ratings-fraud detection," in *Proc. SIAM Int. Conf. Data Mining*, 2016, pp. 495–503.
- [3] S. Ahmed, K. Hinkelmann, and F. Corradini, "Combining machine learning with knowledge engineering to detect fake news in social networks—a survey," *Proc. AAAI Conf. Artif. Intell.*, vol. 12, p. 8, 2019.
- [4] V. Nguyen, K. Sugiyama, P. Nakov, and M. Kan, "Fang: Leveraging social context for fake news detection using graph representation," in *Proc. ACM 29th Int. Conf. Inf. Knowl. Manage.*, 2020, pp. 1165–1174.
- [5] N. T. Tam, M. Weidlich, B. Zheng, H. Yin, N. Q. V. Hung, and B. Stantic, "From anomaly detection to rumour detection using data streams of social platforms," *Very Large DataBases J.*, vol. 12, no. 9, pp. 1016–1029, 2019.
- [6] R. Yu, H. Qiu, Z. Wen, C. Lin, and Y. Liu, "A survey on social media anomaly detection," *SIGKDD Explor.*, vol. 18, no. 1, pp. 1–14, 2016.
- [7] A. Benamira, B. Devillers, E. Lesot, A. K. Ray, M. Saadi, and F. D. Malliaros, "Semi-supervised learning and graph neural networks for fake news detection," in *Proc. Int. Conf. Adv. Social Netw. Anal. Mining*, 2019, pp. 568–569.
- [8] S. Kumar, B. Hooi, D. Makhija, M. Kumar, C. Faloutsos, and V. S. Subrahmanian, "Rev2: Fraudulent user prediction in rating platforms," in *Proc. ACM Int. Conf. Web Search Data Mining*, 2018, pp. 333–341.
- [9] P. Bogdanov, C. Faloutsos, M. Mongiovi, E. E. Papalexakis, R. Ranca, and A. K. Singh, "Netspot: Spotting significant anomalous regions on dynamic networks," in *Proc. SIAM Int. Conf. Data Mining*, 2013, pp. 28–36.
- [10] B. A. Miller, N. Arcolano, and N. T. Bliss, "Efficient anomaly detection in dynamic, attributed graphs: Emerging phenomena and big data," in *Proc. IEEE Int. Conf. Intell. Secur. Inf.*, 2013, pp. 179–184.
- [11] K. Miao, X. Shi, and W. Zhang, "Attack signal estimation for intrusion detection in industrial control system," *Comput. Secur.*, vol. 96, 2020, Art. no. 101926.
- [12] B. Perozzi and L. Akoglu, "Scalable anomaly ranking of attributed neighborhoods," in *Proc. SIAM Int. Conf. Data Mining*, 2016, pp. 207–215.
- [13] X. Luo *et al.*, "ComGA: Community-Aware attributed graph anomaly detection," in *Proc. ACM Int. Conf. Web Search Data Mining*, 2022.
- [14] B. Branco, P. Abreu, A. S. Gomes, M. S. C. Almeida, J. T. Ascensão, and P. Bizarro, "Interleaved sequence RNNS for fraud detection," in *Proc. ACM SIGKDD 26th Int. Conf. Knowl. Discov. Data Mining*, 2020, pp. 3101–3109.
- [15] C. Liu *et al.*, "Fraud transactions detection via behavior tree with local intention calibration," in *Proc. ACM SIGKDD 26th Int. Conf. Knowl. Discov. Data Mining*, 2020, pp. 3035–3043.
- [16] Q. Guo *et al.*, "Securing the deep fraud detector in large-scale e-commerce platform via adversarial machine learning approach," in *Proc. Int. Conf. World Wide Web*, 2019, pp. 616–626.
- [17] B. Igglewick and D. C. Hoaglin, *How to Detect and Handle Outliers*. Milwaukee, WI, USA: ASQ Press, 1993.
- [18] V. Chandola, A. Banerjee, and V. Kumar, "Anomaly detection: A survey," *ACM Comput. Surv.*, vol. 41, no. 3, pp. 15:1–15:58, 2009.
- [19] T. Pourhabibi, O. Kok-Leong, B. H. Kam, and B. Y. Ling, "Fraud detection: A systematic literature review of graph-based anomaly detection approaches," *Decis. Support Syst.*, 2020, Art. no. 113303.
- [20] X. Sun *et al.*, "Detecting users' anomalous emotion using social media for business intelligence," *J. Comput. Sci.*, vol. 25, pp. 193–200, 2018.
- [21] J. Wu, Z. Hong, S. Pan, X. Zhu, Z. Cai, and C. Zhang, "Multi-graph-view learning for graph classification," in *Proc. IEEE Int. Conf. Data Mining*, 2014, pp. 590–599.
- [22] Z. Wang and C. Lan, "Towards a hierarchical bayesian model of multi-view anomaly detection," in *Proc. 29th Int. Joint Conf. Artif. Intell.*, 2020, pp. 2420–2426.
- [23] G. Pang, L. Cao, L. Chen, and H. Liu, "Learning representations of ultrahigh-dimensional data for random distance-based outlier detection," in *Proc. ACM SIGKDD 24th Int. Conf. Knowl. Discov. Data Mining*, 2018, pp. 2041–2050.
- [24] G. Pang *et al.*, "Deep learning for anomaly detection," *ACM Comput. Surv.*, vol. 54, no. 2, pp. 38:1–38:38, 2021.
- [25] L. Akoglu, H. Tong, and D. Koutra, "Graph based anomaly detection and description: A survey," *Data Min. Knowl. Discovery*, vol. 29, no. 3, pp. 626–688, 2015.
- [26] B. Hooi, K. Shin, H. A. Song, A. Beutel, N. Shah, and C. Faloutsos, "Graph-based fraud detection in the face of camouflage," *ACM Trans. Knowl. Discov. Data*, vol. 11, no. 4, pp. 44:1–44:26, 2017.

- [27] Y. Dou, Z. Liu, L. Sun, Y. Deng, H. Peng, and P. S. Yu, "Enhancing graph neural network-based fraud detectors against camouflaged fraudsters," in *Proc. ACM Int. Conf. Inf. Knowl. Manag.*, 2020, pp. 315–324.
- [28] S. Pandit, D. H. Chau, S. Wang, and C. Faloutsos, "Netprobe: Fast and scalable system for fraud detection in online auction networks," in *Proc. 16th Int. Conf. World Wide Web*, 2007, pp. 201–210.
- [29] K. Shin, B. Hooi, J. Kim, and C. Faloutsos, "Densealert: Incremental dense-subtensor detection in tensor streams," in *Proc. ACM SIGKDD 23rd Int. Conf. Knowl. Discov. Data Mining*, 2017, pp. 1057–1066.
- [30] Z. Liu, J. X. Yu, Y. Ke, X. Lin, and L. Chen, "Spotting significant changing subgraphs in evolving graphs," in *Proc. IEEE 8th Int. Conf. Data Mining*, 2008, pp. 917–922.
- [31] J. Wu, Z. Hong, S. Pan, X. Zhu, C. Zhang, and Z. Cai, "Multi-graph learning with positive and unlabeled bags," in *Proc. SIAM Int. Conf. Data Mining*, 2014, pp. 217–225.
- [32] L. Gao *et al.*, "Collaborative dynamic sparse topic regression with user profile evolution for item recommendation," in *Proc. AAAI Conf. Artif. Intell.*, vol. 31, no. 1, 2017, pp. 1316–1322.
- [33] C. C. Aggarwal, Y. Zhao, and S. Y. Philip, "Outlier detection in graph streams," in *Proc. IEEE 27th Int. Conf. Data Eng.*, 2011, pp. 399–409.
- [34] J. Wu, S. Pan, X. Zhu, C. Zhang, and P. S. Yu, "Multiple structure-view learning for graph classification," *IEEE Trans. Neural Netw. Learn. Syst.*, vol. 29, no. 7, pp. 3236–3251, Jul. 2018.
- [35] H. Wang, C. Zhou, X. Chen, J. Wu, S. Pan, and J. Wang, "Graph stochastic neural networks for semi-supervised learning," in *Proc. Int. Conf. Neural Inf. Process. Syst.*, 2020, pp. 1–10.
- [36] Z. Chen, W. Hendrix, and N. F. Samatova, "Community-based anomaly detection in evolutionary networks," *J. Intell. Inf. Syst.*, vol. 39, no. 1, pp. 59–85, 2012.
- [37] F. Jie, C. Wang, F. Chen, L. Li, and X. Wu, "Block-structured optimization for anomalous pattern detection in interdependent networks," in *Proc. IEEE Int. Conf. Data Mining*, 2019, pp. 1138–1143.
- [38] L. Akoglu, M. McGlohon, and C. Faloutsos, "Oddball: Spotting anomalies in weighted graphs," in *Proc. Pacific-Asia Conf. Knowl. Discov. Data Mining*, 2010, pp. 410–421.
- [39] D. Eswaran, C. Faloutsos, S. Guha, and N. Mishra, "Spotlight: Detecting anomalies in streaming graphs," in *Proc. ACM SIGKDD 24th Int. Conf. Knowl. Discov. Data Mining*, 2018, pp. 1378–1386.
- [40] N. Li, H. Sun, K. C. Chipman, J. George, and X. Yan, "A probabilistic approach to uncovering attributed graph anomalies," in *Proc. SIAM Int. Conf. Data Mining*, 2014, pp. 82–90.
- [41] J. Li, H. Dani, X. Hu, and H. Liu, "Radar: Residual analysis for anomaly detection in attributed networks," in *Proc. 26th Int. Joint Conf. Artif. Intell.*, 2017, pp. 2152–2158.
- [42] M. W. Mahoney and P. Drineas, "CUR matrix decompositions for improved data analysis," *Proc. Nat. Acad. Sci. USA*, vol. 106, no. 3, pp. 697–702, 2009.
- [43] S. M. Erfani, S. Rajasegarar, S. Karunasekera, and C. Leckie, "High-dimensional and large-scale anomaly detection using a linear one-class SVM with deep learning," *Pattern Recognit.*, vol. 58, pp. 121–134, 2016.
- [44] S. Thudumu, P. Branch, J. Jin, and J. J. Singh, "A comprehensive survey of anomaly detection techniques for high dimensional big data," *J. Big Data*, vol. 7, no. 1, 2020, Art. no. 42.
- [45] A. Boukerche, L. Zheng, and O. Alfandi, "Outlier detection: Methods, models, and classification," *ACM Comput. Surv.*, vol. 53, no. 3, pp. 55:1–55:37, 2020.
- [46] S. Bulusu, B. Kailkhura, B. Li, P. K. Varshney, and D. Song, "Anomalous instance detection in deep learning: A survey," 2020, *arXiv:2003.06979*.
- [47] R. Chalapathy and S. Chawla, "Deep learning for anomaly detection: A survey," 2019, *arXiv:1901.03407*.
- [48] S. Ranshous, S. Shen, D. Koutra, S. Harenberg, C. Faloutsos, and N. F. Samatova, "Anomaly detection in dynamic networks: A survey," *Wiley Interdiscip. Rev. Comput. Stat.*, vol. 7, no. 3, pp. 223–247, 2015.
- [49] D. J. D'Souza and K. U. K. Reddy, "Anomaly detection for big data using efficient techniques: A review," in *Proc. Adv. Artif. Intell. Data Eng.*, 2021, pp. 1067–1080.
- [50] S. Eltanbouly, M. Bashendy, N. AlNaimi, Z. Chkirkene, and A. Erbad, "Machine learning techniques for network anomaly detection: A survey," in *Proc. Int. Conf. Inform. IoT Enabling Technol.*, 2020, pp. 156–162.
- [51] G. Fernandes, J. J. Rodrigues, L. F. Carvalho, J. F. Al-Muhktadi, and M. L. Proen  a, "A comprehensive survey on network anomaly detection," *Telecommun. Syst.*, vol. 70, no. 3, pp. 447–489, 2019.
- [52] D. Kwon, H. Kim, J. Kim, S. C. Suh, I. Kim, and K. J. Kim, "A survey of deep learning-based network anomaly detection," *Clust. Comput.*, vol. 22, pp. 949–961, 2019.
- [53] P. Gogoi, D. K. Bhattacharyya, B. Borah, and J. K. Kalita, "A survey of outlier detection methods in network anomaly identification," *Comput. J.*, vol. 54, no. 4, pp. 570–588, 2011.
- [54] D. Savage, X. Zhang, X. Yu, P. Chou, and Q. Wang, "Anomaly detection in online social networks," *Soc. Netw.*, vol. 39, pp. 62–70, 2014.
- [55] R. Wang, K. Nie, T. Wang, Y. Yang, and B. Long, "Deep learning for anomaly detection," in *Proc. ACM 13th Int. Conf. Web Search Data Mining*, 2020, pp. 894–896.
- [56] Y. Zhang, J. Wu, Z. Cai, B. Du, and S. Y. Philip, "An unsupervised parameter learning model for RVFL neural network," *Neural Netw.*, vol. 112, pp. 85–97, 2019.
- [57] Q. Wang, W. Zhao, J. Yang, J. Wu, C. Zhou, and Q. Xing, "Atnet-trust: Attributed trust network embedding for trust prediction in online social networks," in *Proc. IEEE Int. Conf. Data Mining*, 2020, pp. 601–610.
- [58] F. Liu *et al.*, "Deep learning for community detection: progress, challenges and opportunities," in *Proc. 29th Int. Joint Conf. Artif. Intell.*, 2020, pp. 4981–4987.
- [59] Z. Wu, S. Pan, F. Chen, G. Long, C. Zhang, and P. S. Yu, "A comprehensive survey on graph neural networks," *IEEE Trans. Neural Netw. Learn. Syst.*, vol. 32, no. 1, pp. 4–24, Jan. 2021.
- [60] P. Cui, X. Wang, J. Pei, and W. Zhu, "A survey on network embedding," *IEEE Trans. Knowl. Data Eng.*, vol. 31, no. 5, pp. 833–852, May 2019.
- [61] S. Zhu, S. Pan, C. Zhou, J. Wu, Y. Cao, and B. Wang, "Graph geometry interaction learning," in *Proc. Int. Conf. Neural Inf. Process. Syst.*, 2020.
- [62] X. Su *et al.*, "A comprehensive survey on community detection with deep learning," 2021, *arXiv:2105.12584*.
- [63] C. C. Noble and D. J. Cook, "Graph-based anomaly detection," in *Proc. ACM SIGKDD 9th Int. Conf. Knowl. Discov. Data Mining*, 2003, pp. 631–636.
- [64] X. Teng, M. Yan, A. M. Ertugrul, and Y. Lin, "Deep into hypersphere: Robust and unsupervised anomaly discovery in dynamic networks," in *Proc. 27th Int. Joint Conf. Artif. Intell.*, 2018, pp. 2724–2730.
- [65] N. Shah *et al.*, "Edgecentric: Anomaly detection in edge-attributed networks," in *Proc. IEEE 16th Int. Conf. Data Mining*, 2016, pp. 327–334.
- [66] B. Wang, N. Z. Gong, and H. Fu, "Gang: Detecting fraudulent users in online social networks via guilt-by-association on directed graphs," in *Proc. IEEE Int. Conf. Data Mining*, 2017, pp. 465–474.
- [67] T. Zhao, C. Deng, K. Yu, T. Jiang, D. Wang, and M. Jiang, "Error-bounded graph anomaly loss for GNNS," in *Proc. ACM Int. Conf. Inf. Knowl. Manage.*, 2020, pp. 1873–1882.
- [68] Q. Zhang and S. Zhu, "Visual interpretability for deep learning: A survey," *Front. Inf. Technol. Electron. Eng.*, vol. 19, no. 1, pp. 27–39, 2018.
- [69] L. Akoglu, "Anomaly mining—past, present and future," in *Proc. ACM 30th Int. Conf. Inf. Knowl. Manage.*, 2021, pp. 1–2.
- [70] Y. Zhao, R. A. Rossi, and L. Akoglu, "Automating outlier detection via meta-learning," 2020, *arXiv:2009.10606*.
- [71] L. Ruf *et al.*, "A unifying review of deep and shallow anomaly detection," *Proc. IEEE*, vol. 109, no. 5, pp. 756–795, May 2021.
- [72] A. Bojchevski and S. G  nnemann, "Bayesian robust attributed graph clustering: Joint learning of partial anomalies and group structure," in *Proc. AAAI 32nd Conf. Artif. Intell.*, 2018, pp. 2738–2745.
- [73] M. Zhu and H. Zhu, "Mixedad: A scalable algorithm for detecting mixed anomalies in attributed graphs," in *Proc. AAAI Conf. Artif. Intell.*, 2020, pp. 1274–1281.
- [74] B. Perozzi, L. Akoglu, P. I. S  nchez, and E. M  ller, "Focused clustering and outlier detection in large attributed graphs," in *Proc. ACM SIGKDD 20th Int. Conf. Knowl. Discov. Data Mining*, 2014, pp. 1346–1355.
- [75] Q. Ding, N. Katenka, P. Barford, E. D. Kolaczyk, and M. Crovella, "Intrusion as (anti)social communication: characterization and detection," in *Proc. ACM SIGKDD 18th Int. Conf. Knowl. Discov. Data Mining*, 2012, pp. 886–894.

- [76] B. Hooi, H. A. Song, A. Beutel, N. Shah, K. Shin, and C. Faloutsos, "Fraudar: Bounding graph fraud in the face of camouflage," in *Proc. ACM SIGKDD 22nd Int. Conf. Knowl. Discov. Data Mining*, 2016, pp. 895–904.
- [77] R. Hu, C. C. Aggarwal, S. Ma, and J. Huai, "An embedding approach to anomaly detection," in *Proc. IEEE 32nd Int. Conf. Data Eng.*, 2016, pp. 385–396.
- [78] G. Karypis and V. Kumar, "A fast and high quality multilevel scheme for partitioning irregular graphs," *J. Sci. Comput.*, vol. 20, no. 1, pp. 359–392, 1998.
- [79] B. Perozzi, R. Al-Rfou, and S. Skiena, "Deepwalk: Online learning of social representations," in *Proc. ACM SIGKDD 20th Int. Conf. Knowl. Discov. Data Mining*, 2014, pp. 701–710.
- [80] A. Grover and J. Leskovec, "Node2vec: Scalable feature learning for networks," in *Proc. ACM SIGKDD 22nd Int. Conf. Knowl. Discov. Data Mining*, 2016, pp. 855–864.
- [81] J. Tang, M. Qu, M. Wang, M. Zhang, J. Yan, and Q. Mei, "Line: Large-scale information network embedding," in *Proc. 24th Int. Conf. World Wide Web*, 2015, pp. 1067–1077.
- [82] S. Bandyopadhyay, L. N. S. V. Vivek, and M. N. Murty, "Outlier resistant unsupervised deep architectures for attributed network embedding," in *Proc. ACM 13th Int. Conf. Web Search Data Mining*, 2020, pp. 25–33.
- [83] S. Bandyopadhyay, N. Lokesh, and M. N. Murty, "Outlier aware network embedding for attributed networks," in *Proc. AAAI Conf. Artif. Intell.*, 2019, pp. 12–19.
- [84] W. Yu, W. Cheng, C. C. Aggarwal, K. Zhang, H. Chen, and W. Wang, "Netwalk: A flexible deep embedding approach for anomaly detection in dynamic networks," in *Proc. ACM SIGKDD 24th Int. Conf. Knowl. Discov. Data Mining*, 2018, pp. 2672–2681.
- [85] L. Cai *et al.*, "Structural temporal graph neural networks for anomaly detection in dynamic graphs," in *Proc. ACM 30th Int. Conf. Inf. Knowl. Manage.*, 2021, pp. 3747–3756.
- [86] M. M. Breunig, H.-P. Kriegel, R. T. Ng, and J. Sander, "LOF: Identifying density-based local outliers," in *Proc. ACM SIGMOD Int. Conf. Manage. Data*, 2000, pp. 93–104.
- [87] C. C. Aggarwal and P. S. Yu, "Outlier detection for high dimensional data," in *Proc. ACM SIGMOD Int. Conf. Manage. Data*, 2001, pp. 37–46.
- [88] W. L. Hamilton, Z. Ying, and J. Leskovec, "Inductive representation learning on large graphs," in *Proc. 31st Int. Conf. Neural Inf. Process. Syst.*, 2017, pp. 1024–1034.
- [89] W. L. Hamilton, R. Ying, and J. Leskovec, "Representation learning on graphs: Methods and applications," *IEEE Data Eng. Bull.*, vol. 40, no. 3, pp. 52–74, Mar. 2017.
- [90] P. Velickovic *et al.*, "Graph attention networks," in *Proc. Int. Conf. Learn. Represent.*, 2018.
- [91] T. N. Kipf and M. Welling, "Semi-supervised classification with graph convolutional networks," in *Proc. Int. Conf. Learn. Represent.*, 2017.
- [92] K. Ding, J. Li, R. Bhanushali, and H. Liu, "Deep anomaly detection on attributed networks," in *Proc. SIAM Int. Conf. Data Mining*, 2019, pp. 594–602.
- [93] Z. Peng, M. Luo, J. Li, L. Xue, and Q. Zheng, "A deep multi-view framework for anomaly detection on attributed networks," *IEEE Trans. Knowl. Data Eng.*, early access, Aug. 07, 2020, doi: [10.1109/TKDE.2020.3015098](https://doi.org/10.1109/TKDE.2020.3015098).
- [94] X.-R. Sheng, D.-C. Zhan, S. Lu, and Y. Jiang, "Multi-view anomaly detection neighborhood in locality matters," in *Proc. AAAI Conf. Artif. Intell.*, 2019, pp. 4894–4901.
- [95] J. Wu, S. Pan, X. Zhu, and Z. Cai, "Boosting for multi-graph classification," *IEEE Trans. Cybern.*, vol. 45, no. 3, pp. 416–429, Mar. 2015.
- [96] J. Wu, X. Zhu, C. Zhang, and Z. Cai, "Multi-instance multi-graph dual embedding learning," in *Proc. IEEE 13th Int. Conf. Data Mining*, 2013, pp. 827–836.
- [97] Y. Li, X. Huang, J. Li, M. Du, and N. Zou, "SPECAE: Spectral autoencoder for anomaly detection in attributed networks," in *Proc. ACM 28th Int. Conf. Inf. Knowl. Manage.*, 2019, pp. 2233–2236.
- [98] J. Wang, R. Wen, C. Wu, Y. Huang, and J. Xion, "FdGars: Fraudster detection via graph convolutional networks in online app review system," in *Proc. Int. Conf. World Wide Web*, 2019, pp. 310–316.
- [99] S. Zhang, H. Yin, T. Chen, Q. V. H. Nguyen, Z. Huang, and L. Cui, "GCN-based user representation learning for unifying robust recommendation and fraudster detection," in *Proc. 43rd Int. ACM SIGIR Conf. Res. Develop. Inf. Retrieval*, 2020, pp. 689–698.
- [100] K. Ding, J. Li, and H. Liu, "Interactive anomaly detection on attributed networks," in *Proc. ACM 12th Int. Conf. Web Search Data Mining*, 2019, pp. 357–365.
- [101] J. Langford and T. Zhang, "The epoch-greedy algorithm for multi-armed bandits with side information," in *Proc. Int. Conf. Neural Inf. Process. Syst.*, 2008, pp. 817–824.
- [102] P. Morales, R. S. Caceres, and T. Eliassi-Rad, "Selective network discovery via deep reinforcement learning on embedded spaces," *Appl. Netw. Sci.*, vol. 6, no. 1, pp. 1–20, 2021.
- [103] R. A. Rossi, B. Gallagher, J. Neville, and K. Henderson, "Modeling dynamic behavior in large evolving graphs," in *Proc. ACM 6th Int. Conf. Web Search Data Mining*, 2013, pp. 667–676.
- [104] H. Wang, J. Wu, W. Hu, and X. Wu, "Detecting and assessing anomalous evolutionary behaviors of nodes in evolving social networks," *ACM Trans. Knowl. Discov. Data*, vol. 13, no. 1, pp. 12:1–12:24, 2019.
- [105] Y. Wang, J. Zhang, S. Guo, H. Yin, C. Li, and H. Chen, "Decoupling representation learning and classification for GNN-based anomaly detection," in *Proc. Int. ACM SIGIR 44th Conf. Res. Develop. Inf. Retrieval*, 2021, pp. 1239–1248.
- [106] N. Liu, X. Huang, and X. Hu, "Accelerated local anomaly detection via resolving attributed networks," in *Proc. 26th Int. Joint Conf. Artif. Intell.*, 2017, pp. 2337–2343.
- [107] Z. Peng, M. Luo, J. Li, H. Liu, and Q. Zheng, "Anomalous: A joint modeling approach for anomaly detection on attributed networks," in *Proc. 27th Int. Joint Conf. Artif. Intell.*, 2018, pp. 3513–3519.
- [108] L. Wu, X. Hu, F. Morstatter, and H. Liu, "Adaptive spammer detection with sparse group modeling," in *Proc. AAAI Int. Conf. Web Soc. Media*, 2017, pp. 319–326.
- [109] Y. Pei, T. Huang, W. van Ipenburg, and M. Pechenizkiy, "ResGCN: Attention-based deep residual modeling for anomaly detection on attributed networks," *Mach. Learn.*, pp. 1–23, 2021.
- [110] H. Fan, F. Zhang, and Z. Li, "Anomalydae: Dual autoencoder for anomaly detection on attributed networks," in *Proc. IEEE Int. Conf. Acoust. Speech Signal Process.*, 2020, pp. 5685–5689.
- [111] D. Wang *et al.*, "A semi-supervised graph attentive network for financial fraud detection," in *Proc. IEEE Int. Conf. Data Mining*, 2019, pp. 598–607.
- [112] K. Ding, J. Li, N. Agarwal, and H. Liu, "Inductive anomaly detection on attributed networks," in *Proc. 29th Int. Joint Conf. Artif. Intell.*, 2020, pp. 1288–1294.
- [113] L. Zhang, J. Yuan, Z. Liu, Y. Pei, and L. Wang, "A robust embedding method for anomaly detection on attributed networks," in *Proc. Int. Joint Conf. Neural Netw.*, 2019, pp. 1–8.
- [114] J. Liang, P. Jacobs, J. Sun, and S. Parthasarathy, "Semi-supervised embedding in attributed networks with outliers," in *Proc. SIAM Int. Conf. Data Mining*, 2018, pp. 153–161.
- [115] X. Wang, B. Jin, Y. Du, P. Cui, Y. Tan, and Y. Yang, "One-class graph neural networks for anomaly detection in attributed networks," *Neural Comput. Appl.*, vol. 33, pp. 12073–12085, 2021.
- [116] Y. Liu, Z. Li, S. Pan, C. Gong, C. Zhou, and G. Karypis, "Anomaly detection on attributed networks via contrastive self-supervised learning," *IEEE Trans. Neural Netw. Learn. Syst.*, early access, Apr. 05, 2021, doi: [10.1109/TNNLS.2021.3068344](https://doi.org/10.1109/TNNLS.2021.3068344).
- [117] K. Ding, K. Shu, X. Shan, J. Li, and H. Liu, "Cross-domain graph anomaly detection," *IEEE Trans. Neural Netw. Learn. Syst.*, early access, Oct. 01, 2021, doi: [10.1109/TNNLS.2021.3110982](https://doi.org/10.1109/TNNLS.2021.3110982).
- [118] G. Zhang *et al.*, "Fraudre: Fraud detection dual-resistant to graph inconsistency and imbalance," in *Proc. IEEE Int. Conf. Data Mining*, 2021.
- [119] K. Ding, Q. Zhou, H. Tong, and H. Liu, "Few-shot network anomaly detection via cross-network meta-learning," in *Proc. Web Conf.*, 2021, pp. 2448–2456.
- [120] X. Teng, Y. Lin, and X. Wen, "Anomaly detection in dynamic networks using multi-view time-series hypersphere learning," in *Proc. ACM Int. Conf. Inf. Knowl. Manage.*, 2017, pp. 827–836.
- [121] P. Zheng, S. Yuan, X. Wu, J. Li, and A. Lu, "One-class adversarial nets for fraud detection," in *Proc. AAAI Conf. Artif. Intell.*, 2019, pp. 1286–1293.
- [122] N. Ailon *et al.*, "Streaming k-means approximation," in *Proc. Int. Conf. Neural Inf. Process. Syst.*, 2009, pp. 10–18.
- [123] I. J. Goodfellow *et al.*, "Generative adversarial nets," in *Proc. Int. Conf. Neural Inf. Process. Syst.*, 2014, pp. 2672–2680.
- [124] Z. Dai *et al.*, "Good semi-supervised learning that requires a bad GAN," in *Proc. Int. Conf. Neural Inf. Process. Syst.*, 2017, pp. 6510–6520.
- [125] N. Srivastava, E. Mansimov, and R. Salakhutdinov, "Unsupervised learning of video representations using LSTMS," in *Proc. Int. Conf. Mach. Learn.*, 2015, pp. 843–852.

- [126] Y.-Y. Chang, P. Li, R. Sosic, M. Afifi, M. Schweighauser, and J. Leskovec, "F-Fade: Frequency factorization for anomaly detection in edge streams," in *Proc. ACM 14th Int. Conf. Web Search Data Mining*, 2021, pp. 589–597.
- [127] L. Ouyang, Y. Zhang, and Y. Wang, "Unified graph embedding-based anomalous edge detection," in *Proc. Int. Joint Conf. Neural Netw.*, 2020, pp. 1–8.
- [128] D. Duan, L. Tong, Y. Li, J. Lu, L. Shi, and C. Zhang, "Aane: Anomaly aware network embedding for anomalous link detection," in *Proc. IEEE Int. Conf. Data Mining*, 2020, pp. 1002–1007.
- [129] L. Xu, X. Wei, J. Cao, and P. S. Yu, "ICANE: Interaction content-aware network embedding via co-embedding of nodes and edges," *Int. J. Data Sci. Anal.*, vol. 9, no. 4, pp. 401–414, 2020.
- [130] S. Ranshous, S. Harenberg, K. Sharma, and N. F. Samatova, "A scalable approach for outlier detection in edge streams using sketch-based approximations," in *Proc. SIAM Int. Conf. Data Mining*, 2016, pp. 189–197.
- [131] L. Zheng, Z. Li, J. Li, Z. Li, and J. Gao, "Addgraph: Anomaly detection in dynamic graph using attention-based temporal GCN," in *Proc. Int. Joint Conf. Artif. Intell.*, 2019, pp. 4419–4425.
- [132] Q. Cui, S. Wu, Y. Huang, and L. Wang, "A hierarchical contextual attention-based network for sequential recommendation," *Neurocomputing*, vol. 358, pp. 141–149, 2019.
- [133] M. Shao, J. Li, F. Chen, and X. Chen, "An efficient framework for detecting evolving anomalous subgraphs in dynamic networks," in *Proc. IEEE Conf. Comput. Commun.*, 2018, pp. 2258–2266.
- [134] H. Wang, C. Zhou, J. Wu, W. Dang, X. Zhu, and J. Wang, "Deep structure learning for fraud detection," in *Proc. IEEE Int. Conf. Data Mining*, 2018, pp. 567–576.
- [135] M. Ester, H. Kriegel, J. Sander, and X. Xu, "A density-based algorithm for discovering clusters in large spatial databases with noise," in *Proc. ACM Int. Conf. Knowl. Discov. Data Mining*, 1996, pp. 226–231.
- [136] M. Zheng, C. Zhou, J. Wu, S. Pan, J. Shi, and L. Guo, "Fraudne: A joint embedding approach for fraud detection," in *Proc. Int. Joint Conf. Neural Netw.*, 2018, pp. 1–8.
- [137] J. Wu, X. Zhu, C. Zhang, and P. S. Yu, "Bag constrained structure pattern mining for multi-graph classification," *IEEE Trans. Knowl. Data Eng.*, vol. 26, no. 10, pp. 2382–2396, Oct. 2014.
- [138] Q. Sun *et al.*, "Sugar: Subgraph neural network with reinforcement pooling and self-supervised mutual information mechanism," in *Proc. Int. Conf. World Wide Web*, 2021, pp. 2081–2091.
- [139] E. A. Manzoor, S. M. Milajerdi, and L. Akoglu, "Fast memory-efficient anomaly detection in streaming heterogeneous graphs," in *Proc. ACM SIGKDD 22nd Int. Conf. Knowl. Discov. Data Mining*, 2016, pp. 1035–1044.
- [140] B. Hooi *et al.*, "Changedar: Online localized change detection for sensor data on a graph," in *Proc. ACM 27th Int. Conf. Inf. Knowl. Manage.*, 2018, pp. 507–516.
- [141] Y. Dou *et al.*, "User preference-aware fake news detection," in *Proc. 44th Int. ACM SIGIR Conf. Res. Develop. Inf. Retrieval*, 2021, pp. 2051–2055.
- [142] L. Zhao and L. Akoglu, "On using classification datasets to evaluate graph outlier detection: Peculiar observations and new insights," 2020, *arXiv:2012.12931*.
- [143] L. Ruf *et al.*, "Deep one-class classification," in *Proc. Int. Conf. Mach. Learn.*, 2018, pp. 4393–4402.
- [144] A. Narayanan, M. Chandramohan, R. Venkatesan, L. Chen, Y. Liu, and S. Jaiswal, "Graph2vec: Learning distributed representations of graphs," 2017, *arXiv:1707.05005*.
- [145] S. Verma and Z.-L. Zhang, "Hunt for the unique, stable, sparse and fast feature learning on graphs," in *Proc. 31st Int. Conf. Neural Inf. Process. Syst.*, 2017, pp. 87–97.
- [146] G. Zhang *et al.*, "efraudcom: An e-commerce fraud detection system via competitive graph neural networks," *ACM Trans. Inf. Syst.*, 2021.
- [147] X. Teng, M. Yan, A. M. Ertugrul, and Y. Lin, "Deep into hypersphere: Robust and unsupervised anomaly discovery in dynamic networks," in *Proc. 27th Int. Joint Conf. Artif. Intell.*, 2018, pp. 2724–2730.
- [148] Y. Wan, Y. Liu, D. Wang, and Y. Wen, "Glad-paw: Graph-based log anomaly detection by position aware weighted graph attention network," in *Pacific-Asia Conf. Knowl. Discov. Data Mining*, 2021, pp. 66–77.
- [149] Z. Zhang, J. Jia, B. Wang, and N. Z. Gong, "Backdoor attacks to graph neural networks," in *Proc. 26th ACM Symp. Access Control Models Technol.*, 2021, pp. 15–26.
- [150] H. Dai *et al.*, "Adversarial attack on graph structured data," in *Proc. Int. Conf. Mach. Learn.*, 2018, pp. 1115–1124.
- [151] X. Lin *et al.*, "Exploratory adversarial attacks on graph neural networks," in *Proc. IEEE Int. Conf. Data Mining*, 2020, pp. 1136–1141.
- [152] D. Eswaran and C. Faloutsos, "Sedanspot: Detecting anomalies in edge streams," in *Proc. IEEE Int. Conf. Data Mining*, 2018, pp. 953–958.
- [153] L. Gutiérrez-Gómez, A. Bovet, and J.-C. Delvenne, "Multi-scale anomaly detection on attributed networks," in *Proc. AAAI Conf. Artif. Intell.*, 2020, pp. 678–685.
- [154] H. Nilforoshan and N. Shah, "Slicendice: Mining suspicious multi-attribute entity groups with multi-view graphs," in *Proc. IEEE Int. Conf. Data Sci. Adv. Anal.*, 2019, pp. 351–363.
- [155] Z. Liu, Y. Dou, P. S. Yu, Y. Deng, and H. Peng, "Alleviating the inconsistency problem of applying graph neural network to fraud detection," in *Proc. 43rd Int. ACM SIGIR Conf. Res. Develop. Inf. Retrieval*, 2020, pp. 1569–1572.
- [156] Y.-J. Lu and C.-T. Li, "Gcan: Graph-aware co-attention networks for explainable fake news detection on social media," in *Proc. Annu. Meet. Assoc. Comput. Linguist. Proc. Conf.*, 2020, pp. 505–514.
- [157] Q. Huang, J. Yu, J. Wu, and B. Wang, "Heterogeneous graph attention networks for early detection of rumors on twitter," in *Proc. Int. Joint Conf. Neural Netw.*, 2020, pp. 1–8.
- [158] C. Yuan, Q. Ma, W. Zhou, J. Han, and S. Hu, "Jointly embedding the local and global relations of heterogeneous graph for rumor detection," in *Proc. IEEE Int. Conf. Data Mining*, 2019, pp. 796–805.
- [159] M. Yoon, B. Hooi, K. Shin, and C. Faloutsos, "Fast and accurate anomaly detection in dynamic graphs with a two-pronged approach," in *Proc. ACM SIGKDD 25th Int. Conf. Knowl. Discov. Data Mining*, 2019, pp. 647–657.
- [160] B. Zong *et al.*, "Deep autoencoding gaussian mixture model for unsupervised anomaly detection," in *Proc. Int. Conf. Learn. Represent.*, 2018.
- [161] G. Pang, C. Shen, and A. van den Hengel, "Deep anomaly detection with deviation networks," in *Proc. ACM SIGKDD 25th Int. Conf. Knowl. Discov. Data Mining*, 2019, pp. 353–362.
- [162] C. Zhou and R. C. Paffenroth, "Anomaly detection with robust deep autoencoders," in *Proc. ACM SIGKDD 23rd Int. Conf. Knowl. Discov. Data Mining*, 2017, pp. 665–674.
- [163] R. Chalapathy, E. Toth, and S. Chawla, "Group anomaly detection using deep generative models," in *Proc. Eur. Conf. Mach. Learn. Knowl. Discov. Databases*, 2018, pp. 173–189.
- [164] Z. Liu, C. Chen, X. Yang, J. Zhou, X. Li, and L. Song, "Heterogeneous graph neural networks for malicious account detection," in *Proc. ACM 27th Int. Conf. Inf. Knowl. Manage.*, 2018, pp. 2077–2085.
- [165] S. Bhatia, B. Hooi, M. Yoon, K. Shin, and C. Faloutsos, "Midas: Microcluster-based detector of anomalies in edge streams," in *Proc. AAAI Conf. Artif. Intell.*, 2020, pp. 3242–3249.
- [166] S. Dhawan, S. C. R. Gangireddy, S. Kumar, and T. Chakraborty, "Spotting collective behaviour of online frauds in customer reviews," in *Proc. Int. Joint Conf. Artif. Intell.*, 2019, pp. 245–251.
- [167] Q. Huang *et al.*, "Deep spatial-temporal structure learning for rumor detection on twitter," *Neural Comput. Appl.*, pp. 1–11, 2020.
- [168] L. Ruf *et al.*, "Deep semi-supervised anomaly detection," in *Proc. Int. Conf. Learn. Represent.*, 2020.
- [169] S. Kim, *et al.*, "Date: Dual attentive tree-aware embedding for customs fraud detection," in *Proc. ACM SIGKDD 26th Int. Conf. Knowl. Discov. Data Mining*, 2020, pp. 2880–2890.
- [170] S. Rayana and L. Akoglu, "Less is more: Building selective anomaly ensembles," *ACM Trans. Knowl. Discov. Data*, vol. 10, no. 4, pp. 1–33, 2016.
- [171] P. Sen, G. Namata, M. Bilgic, L. Getoor, B. Gallagher, and T. Eliassi-Rad, "Collective classification in network data," *AI Mag.*, vol. 29, no. 3, pp. 93–106, 2008.
- [172] P. I. Sánchez, E. Müller, F. Laforet, F. Keller, and K. Böhm, "Statistical selection of congruent subspaces for mining attributed graphs," in *Proc. IEEE 13th Int. Conf. Data Mining*, 2013, pp. 647–656.
- [173] A. Lancichinetti and S. Fortunato, "Community detection algorithms: A comparative analysis," *Phys. Rev. E*, vol. 80, no. 5, 2009, Art. no. 056117.
- [174] D. J. Watts and S. H. Strogatz, "Collective dynamics of 'small-world' networks," *Nature*, vol. 393, no. 6684, pp. 440–442, 1998.
- [175] L. Akoglu and C. Faloutsos, "RTG: A recursive realistic graph generator using random typing," in *Proc. Eur. Conf. Mach. Learn. Knowl. Discov. Databases*, 2009, pp. 13–28.
- [176] L. Bozarth and C. Budak, "Toward a better performance evaluation framework for fake news classification," in *Proc. AAAI Int. Conf. Web Soc. Media*, 2020, pp. 60–71.

- [177] N. Elmkrab, F. Zhou, F. Li, and H. Zhou, "Evaluation of machine learning algorithms for anomaly detection," in *Proc. Cyber Secur.*, 2020, pp. 1–8.
- [178] A. H. Engly, A. R. Larsen, and W. Meng, "Evaluation of anomaly-based intrusion detection with combined imbalance correction and feature selection," in *Proc. Int. Conf. Netw. Secur.*, 2020, pp. 277–291.
- [179] E. Alsentzer, S. G. Finlayson, M. M. Li, and M. Zitnik, "Subgraph neural networks," in *Proc. Int. Conf. Neural Inf. Process. Syst.*, 2020.
- [180] H. Wang, Q. Zhang, J. Wu, S. Pan, and Y. Chen, "Time series feature learning with labeled and unlabeled data," *Pattern Recognit.*, vol. 89, pp. 55–66, 2019.
- [181] Q. Zhang, J. Wu, P. Zhang, G. Long, and C. Zhang, "Salient subsequence learning for time series clustering," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 41, no. 9, pp. 2193–2207, Sep. 2018.
- [182] A. Sankar, Y. Wu, L. Gou, W. Zhang, and H. Yang, "DYSAT: Deep neural representation learning on dynamic graphs via self-attention networks," in *Proc. ACM Int. Conf. Web Search Data Mining*, 2020, pp. 519–527.
- [183] S. M. Kazemi *et al.*, "Representation learning for dynamic graphs: A survey," *J. Mach. Learn. Res.*, vol. 21, pp. 70:1–70:73, 2020.
- [184] W. L. Hamilton, Z. Ying, and J. Leskovec, "Inductive representation learning on large graphs," in *Proc. 31st Int. Conf. Neural Inf. Process. Syst.*, 2017, pp. 1024–1034.
- [185] M. R. Khan and J. E. Blumenstock, "Multi-GCN: Graph convolutional networks for multi-view networks, with applications to global poverty," in *Proc. AAAI Conf. Artif. Intell.*, vol. 33, no. 01, 2019, pp. 606–613.
- [186] S. Fan, X. Wang, C. Shi, E. Lu, K. Lin, and B. Wang, "One2multi graph autoencoder for multi-view graph clustering," in *Proc. Int. Conf. World Wide Web*, 2020, pp. 3070–3076.
- [187] J. Cheng, Q. Wang, Z. Tao, D. Xie, and Q. Gao, "Multi-view attribute graph convolution networks for clustering," in *Proc. Int. Joint Conf. Artif. Intell.*, 2020, pp. 2973–2979.
- [188] J. Zhang, B. Cao, S. Xie, C. Lu, P. S. Yu, and A. B. Ragin, "Identifying connectivity patterns for brain diseases via multi-side-view guided deep architectures," in *Proc. SIAM Int. Conf. Data Mining*, 2016, pp. 36–44.
- [189] H. Xiao, J. Gao, D. S. Turaga, L. H. Vu, and A. Biem, "Temporal multi-view inconsistency detection for network traffic analysis," in *Proc. 24th Int. Conf. World Wide Web*, 2015, pp. 455–465.
- [190] E. Gujral, R. Pasricha, and E. Papalexakis, "Beyond rank-1: Discovering rich community structure in multi-aspect graphs," in *Proc. Int. Conf. World Wide Web*, 2020, pp. 452–462.
- [191] N. Shah, A. Beutel, B. Gallagher, and C. Faloutsos, "Spotting suspicious link behavior with FBOX: An adversarial perspective," in *Proc. IEEE Int. Conf. Data Mining*, 2014, pp. 959–964.
- [192] H. Chen, H. Yin, X. Sun, T. Chen, B. Gabrys, and K. Musial, "Multi-level graph convolutional networks for cross-platform anchor link prediction," in *Proc. ACM SIGKDD 26th Int. Conf. Knowl. Discov. Data Mining*, 2020, pp. 1503–1511.
- [193] A. Guzzo, A. Pugliese, A. Rullo, D. Saccà, and A. Piccolo, "Malevolent activity detection with hypergraph-based models," *IEEE Trans. Knowl. Data Eng.*, vol. 29, no. 5, pp. 1115–1128, May 2017.
- [194] X. Sun *et al.*, "Heterogeneous hypergraph embedding for graph classification," in *Proc. ACM 14th Int. Conf. Web Search Data Mining*, 2021, pp. 725–733.
- [195] J. Silva and R. Willett, "Hypergraph-based anomaly detection of high-dimensional co-occurrences," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 31, no. 3, pp. 563–569, Mar. 2008.
- [196] J. Tang, C. C. Aggarwal, and H. Liu, "Node classification in signed social networks," in *Proc. SIAM Int. Conf. Data Mining*, 2016, pp. 54–62.
- [197] S. Gao, L. Denoyer, and P. Gallinari, "Temporal link prediction by integrating content and structure information," in *Proc. ACM 20th Int. Conf. Inf. Knowl. Manage.*, 2011, pp. 1169–1174.
- [198] V. Sanh, T. Wolf, and S. Ruder, "A hierarchical multi-task approach for learning embeddings from semantic tasks," in *Proc. AAAI Conf. Artif. Intell.*, vol. 33, no. 01, 2019, pp. 6949–6956.
- [199] M. Hessel, H. Soyer, L. Espeholt, W. Czarnecki, S. Schmitt, and H. van Hasselt, "Multi-task deep reinforcement learning with popart," in *Proc. AAAI Conf. Artif. Intell.*, 2019, pp. 3796–3803.
- [200] G. Pang, C. Yan, C. Shen, A. v. d. Hengel, and X. Bai, "Self-trained deep ordinal regression for end-to-end video anomaly detection," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit.*, 2020, pp. 12173–12182.
- [201] B. Sanchez-Lengeling *et al.*, "Evaluating attribution for graph neural networks," in *Proc. Int. Conf. Neural Inf. Process. Syst.*, 2020.
- [202] T. Idé, A. Dhurandhar, J. Navrátil, M. Singh, and N. Abe, "Anomaly attribution with likelihood compensation," in *Proc. AAAI Conf. Artif. Intell.*, vol. 35, no. 5, 2021, pp. 4131–4138.
- [203] F. Hohman, H. Park, C. Robinson, and D. H. P. Chau, "Summit: Scaling deep learning interpretability by visualizing activation and attribution summarizations," *IEEE Trans. Vis. Comput. Graph.*, vol. 26, no. 1, pp. 1096–1106, 2019.
- [204] M. Kulldorff, "A spatial scan statistic," *Commun. Stat.- Theory Methods*, vol. 26, no. 6, pp. 1481–1496, 1997.
- [205] D. B. Neill, A. W. Moore, M. Sabhnani, and K. Daniel, "Detection of emerging space-time clusters," in *Proc. ACM SIGKDD 11th Int. Conf. Knowl. Discov. Data Mining*, 2005, pp. 218–227.
- [206] J. L. Sharpnack, A. Krishnamurthy, and A. Singh, "Near-optimal anomaly detection in graphs using lovasz extended scan statistic," in *Proc. 26th Int. Conf. Neural Inf. Process. Syst.*, 2013, pp. 1959–1967.
- [207] R. H. Berk and D. H. Jones, "Goodness-of-fit test statistics that dominate the kolmogorov statistics," *Wahrscheinl. Verw. Gebiete.*, vol. 47, no. 1, pp. 47–59, 1979.
- [208] T. Babaie, S. Chawla, S. Ardon, and Y. Yu, "A unified approach to network anomaly detection," in *Proc. BigData*, 2014, pp. 650–655.
- [209] X. Li, P. Chen, L. Jing, Z. He, and G. Yu, "Swisslog: Robust and unified deep learning based log anomaly detection for diverse faults," in *Proc. IEEE 31st Int. Symp. Softw. Reliable Eng.*, 2020, pp. 92–103.



Xiaoxiao Ma received the master's degree from the University of Melbourne, Australia. He is currently working toward the PhD degree with the School of Computing, Macquarie University, Sydney, Australia. His research interests include data mining, deep learning, anomaly detection, machine learning, and social networks.



Jia Wu (Senior Member, IEEE) received the PhD degree from the University of Technology Sydney, Australia. He is currently a senior lecturer and an ARC DECRA fellow with the School of Computing, Macquarie University, Sydney, Australia. He has authored or coauthored more than 100 refereed journal and conference papers, including *IEEE Transactions on Knowledge and Data Engineering*, *IEEE Transactions on Pattern Analysis and Machine Intelligence*, NIPS, KDD, ICDM, and WWW in the areas of his research interests, which include data mining, machine learning, social networks, and graph mining. He was the recipient of SDM'18 Best Paper Award in Data Science Track, IJCNN'17 Best Student Paper Award, and ICDM'14 Best Paper Candidate Award. He is currently an associate editor for *ACM Transactions on Knowledge Discovery from Data*.



Shan Xue received the PhD in computer science from the Center for Artificial Intelligence, University of Technology Sydney, Australia, in 2019 and the PhD in information management from the School of Management, Shanghai University, Shanghai, China, in 2018. She is currently a postdoctoral research fellow with the School of Computing, Faculty of Science and Engineering, Macquarie University, Sydney, Australia and a researcher with CSIRO's Data61, Sydney, Australia. Her research interests include artificial intelligence, machine learning, and knowledge mining.



Jian Yang received the PhD degree in data integration from Australian National University in 1995. She is currently a full professor with the School of Computing, Macquarie University. Prof. Yang has published more than 200 journal and conference papers in international journals and conferences, including *IEEE Transactions on Information Systems, Data and Knowledge Engineering*, VLDB, ICDE, ICDM, and CIKM. Her research interests include business process management, data science, social networks. She is currently on the executive committee for the Computing Research and Education Association of Australasia.



Chuan Zhou received the PhD degree from the Chinese Academy of Sciences in 2013. He is currently an associate professor with the Academy of Mathematics and Systems Science, Chinese Academy of Sciences. He has authored or coauthored more than 80 papers, including *IEEE Transactions on Knowledge and Data Engineering*, ICDM, AAAI, IJCAI, and WWW. His research interests include social network analysis and graph mining. He was the recipient of the outstanding doctoral dissertation of Chinese Academy of Sciences in 2014, the Best Paper Award of ICCS-14, and the Best Student Paper Award of IJCNN-17.



Quan Z. Sheng received the PhD degree in computer science from the University of New South Wales. He is currently a full professor and the head of School of Computing, Macquarie University, Sydney, Australia. He is a postdoc as a research scientist with CSIRO ICT Centre. He has more than 400 publications. His research interests include big data analytics, service oriented computing, knowledge discovery, Internet computing, and the Internet of Things. He was the recipient of AMiner Most Influential Scholar Award in IoT in 2019, ARC Future Fellowship in 2014, Chris Wallace Award for Outstanding Research Contribution in 2012, and Microsoft Fellowship in 2003. He is ranked by Microsoft Academic as one of the Most Impactful Authors in Services Computing (ranked the 4th all time).



Hui Xiong (Fellow, IEEE) received the PhD degree in computer science from the University of Minnesota, USA. He is currently a professor with Rutgers, State University of New Jersey. He was on the organization and program committees of numerous conferences, including KDD, ICDM, and SDM. His research interests include data mining and its business applications. He was elected an ACM distinguished scientist in 2014. He is currently the co-editor-in-chief of *Encyclopedia of GIS* and an associate editor for *IEEE Transactions on Data and Knowledge Engineering*, *IEEE Transactions on Big Data*, *ACM Transactions on Knowledge Discovery from Data*, and *ACM Transactions on Management Information Systems*. He was the recipient of the 2021 AAAI Best Paper Award and the 2011 IEEE ICDM Best Research Paper Award. He is a Fellow of the AAAS.



Leman Akoglu received the PhD degree from Computer Science Department or the School of Computer Science, Carnegie Mellon University, in 2012. She is currently the Heinz College Dean's associate professor of information systems with Carnegie Mellon University. She holds courtesy appointments with (CSD) and Machine Learning Department, School of Computer Science. She also holds three U.S. patents filed by IBM T. J. Watson Research Labs. Her research interests include machine learning and graph mining, focusing on pattern discovery and anomaly mining, with applications to fraud and event detection. Her research has been supported by the NSF, US ARO, DARPA, Adobe, Facebook, Northrop Grumman, PNC Bank, and PwC. She was the recipient of the SDM/IBM Early Career Data Mining Research Award (2020), National Science Foundation CAREER Award (2015), US Army Research Office Young Investigator Award (2013), Eight publication awards, The Most Influential Paper Award (PAKDD 2020), Best Research Paper (SIAM SDM 2019), Best Student Machine Learning Paper Runner-up (ECML PKDD 2018), Best Paper Runner-up (SIAM SDM 2016), Best Research Paper (SIAM SDM 2015), Best Paper (ADC 2014), Best Paper (PAKDD 2010), and Best Knowledge Discovery Paper (ECML PKDD 2009).

▷ For more information on this or any other computing topic, please visit our Digital Library at www.computer.org/csl.