


```
In [6]: train_ds, train_valid_ds = [torchvision.datasets.ImageFolder(
    os.path.join(data_dir, 'train_valid_test', folder),
    transform=transform_train) for folder in ['train', 'train_valid']]

valid_ds, test_ds = [torchvision.datasets.ImageFolder(
    os.path.join(data_dir, 'train_valid_test', folder),
    transform=transform_test) for folder in ['valid', 'test']]
```

```
In [7]: train_iter, train_valid_iter = [torch.utils.data.DataLoader(
    dataset, batch_size, shuffle=True, drop_last=True)
    for dataset in (train_ds, train_valid_ds)]

valid_iter = torch.utils.data.DataLoader(valid_ds, batch_size, shuffle=False,
    drop_last=True)

test_iter = torch.utils.data.DataLoader(test_ds, batch_size, shuffle=False,
    drop_last=False)
```

```
In [8]: def get_net(devices):
    finetune_net = nn.Sequential()
    finetune_net.features = torchvision.models.resnet34(pretrained=True)
    # 定义一个新的输出网络，共有120个输出类别
    finetune_net.output_new = nn.Sequential(nn.Linear(1000, 256),
        nn.ReLU(),
        nn.Linear(256, 120))

    # 将模型参数分配给用于计算的CPU或GPU
    finetune_net = finetune_net.to(devices[0])
    # 冻结参数
    for param in finetune_net.features.parameters():
        param.requires_grad = False
    return finetune_net
```

```
In [9]: loss = nn.CrossEntropyLoss(reduction='none')

def evaluate_loss(data_iter, net, devices):
    l_sum, n = 0.0, 0
    for features, labels in data_iter:
        features, labels = features.to(devices[0]), labels.to(devices[0])
        outputs = net(features)
        l = loss(outputs, labels)
        l_sum += l.sum()
        n += labels.numel()
    return (l_sum / n).to('cpu')
```

```
In [10]: def train(net, train_iter, valid_iter, num_epochs, lr, wd, devices, lr_period,
    lr_decay):
    # 只训练小型自定义输出网络
    net = nn.DataParallel(net, device_ids=devices).to(devices[0])
    trainer = torch.optim.SGD((param for param in net.parameters()
        if param.requires_grad), lr=lr,
        momentum=0.9, weight_decay=wd)
    scheduler = torch.optim.lr_scheduler.StepLR(trainer, lr_period, lr_decay)
    num_batches, timer = len(train_iter), d2l.Timer()
    legend = ['train loss']
    if valid_iter is not None:
        legend.append('valid loss')
    animator = d2l.Animator(xlabel='epoch', xlim=[1, num_epochs],
        legend=legend)
```

```

for epoch in range(num_epochs):
    metric = d2l.Accumulator(2)
    for i, (features, labels) in enumerate(train_iter):
        timer.start()
        features, labels = features.to(devices[0]), labels.to(devices[0])
        trainer.zero_grad()
        output = net(features)
        l = loss(output, labels).sum()
        l.backward()
        trainer.step()
        metric.add(l, labels.shape[0])
        timer.stop()
        if (i + 1) % (num_batches // 5) == 0 or i == num_batches - 1:
            animator.add(epoch + (i + 1) / num_batches,
                          (metric[0] / metric[1], None))
    measures = f'train loss {metric[0] / metric[1]:.3f}'
    if valid_iter is not None:
        valid_loss = evaluate_loss(valid_iter, net, devices)
        animator.add(epoch + 1, (None, valid_loss.detach().cpu()))
    scheduler.step()
    if valid_iter is not None:
        measures += f', valid loss {valid_loss:.3f}'
    print(measures + f'\n{metric[1] * num_epochs / timer.sum():.1f}'
          f' examples/sec on {str(devices)}')

```

```

In [ ]: devices, num_epochs, lr, wd = d2l.try_all_gpus(), 10, 1e-4, 1e-4
        lr_period, lr_decay, net = 2, 0.9, get_net(devices)
        train(net, train_iter, valid_iter, num_epochs, lr, wd, devices, lr_period,
              lr_decay)

```

c:\Users\DELL\anaconda3\envs\pytorch\lib\site-packages\torchvision\models_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.

warnings.warn(

c:\Users\DELL\anaconda3\envs\pytorch\lib\site-packages\torchvision\models_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=ResNet34_Weights.IMAGENET1K_V1`. You can also use `weights=ResNet34_Weights.DEFAULT` to get the most up-to-date weights.

warnings.warn(msg)

Downloading: "https://download.pytorch.org/models/resnet34-b627a593.pth" to C:\Users\DELL\.cache\torch\hub\checkpoints\resnet34-b627a593.pth

```

In [ ]: net = get_net(devices)
        train(net, train_valid_iter, None, num_epochs, lr, wd, devices, lr_period,
              lr_decay)

        preds = []
        for data, label in test_iter:
            output = torch.nn.functional.softmax(net(data.to(devices[0])), dim=1)
            preds.extend(output.cpu().detach().numpy())
        ids = sorted(os.listdir(
            os.path.join(data_dir, 'train_valid_test', 'test', 'unknown')))
        with open('submission.csv', 'w') as f:
            f.write('id,' + ','.join(train_valid_ds.classes) + '\n')
            for i, output in zip(ids, preds):
                f.write(i.split('.')[0] + ',' + ','.join(
                    [str(num) for num in output]) + '\n')

```